



ACDS Lecture Series

Lecture - 10

CSIR

Supervised and Unsupervised Learning

G. N. Sastry and Team

ADVANCED COMPUTATION AND DATA SCIENCES (ACDS) DIVISION

CSIR-North East Institute of Science and Technology, Jorhat, Assam, India

10.1 Introduction:

10.2 Supervised Learning

10.2.1 Classification

10.2.1.1 Random Forest

10.2.1.2 Decision Trees

10.2.1.3 Logistic Regression

10.2.1.4 Support Vector Machines

10.2.1.5 KNN

10.2.1.6 Naïve Bayes

10.2.1.7 Examples

10.2.2 Regression

10.2.2.1 Linear Regression

10.2.2.2 Regularization Techniques (LASSO)

10.2.2.3 Polynomial Regression

10.2.2.4 Examples

10.3 Unsupervised Learning

10.3.1 Clustering

10.3.1.1 K-Means

10.3.1.2 K Nearest Neighbors

10.3.1.3 Association Rule Learning

10.3.2 Dimensionality Reduction

10.3.2.1 PCA

10.3.2.2 LDA

10.3.2.3 GDA

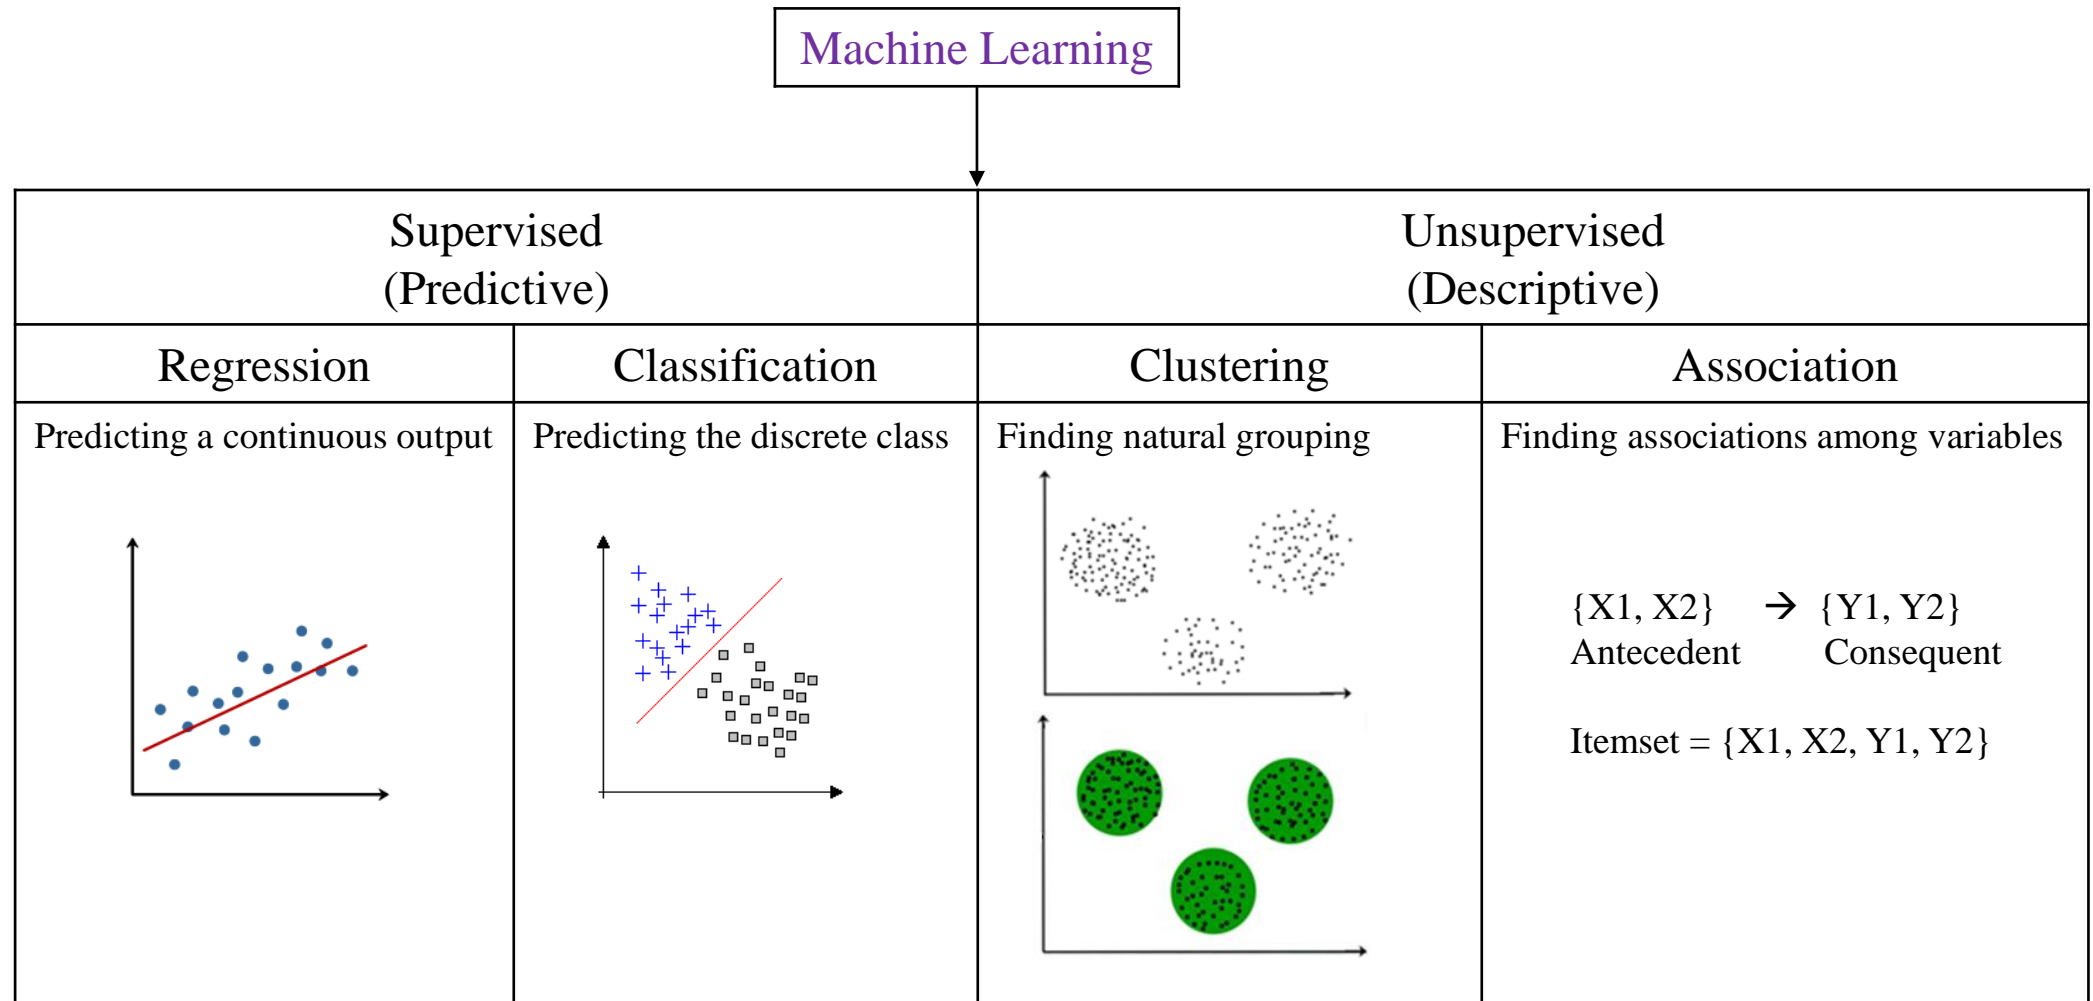
10.3.2.4 Examples

10.4 Exercises

10.5 References

10.5.1 Books

10.5.2 Video Lectures



The Topics Covered in This Section

10.2.1 Classification

10.2.1.1 Random Forest

10.2.1.2 Decision Trees

10.2.1.3 Logistic Regression

10.2.1.4 Support Vector Machines

10.2.1.5 KNN

10.2.1.6 Naïve Bayes

10.2.1.7 Examples

10.2.2 Regression

10.2.2.1 Linear Regression

10.2.2.2 Regularization Techniques (LASSO)

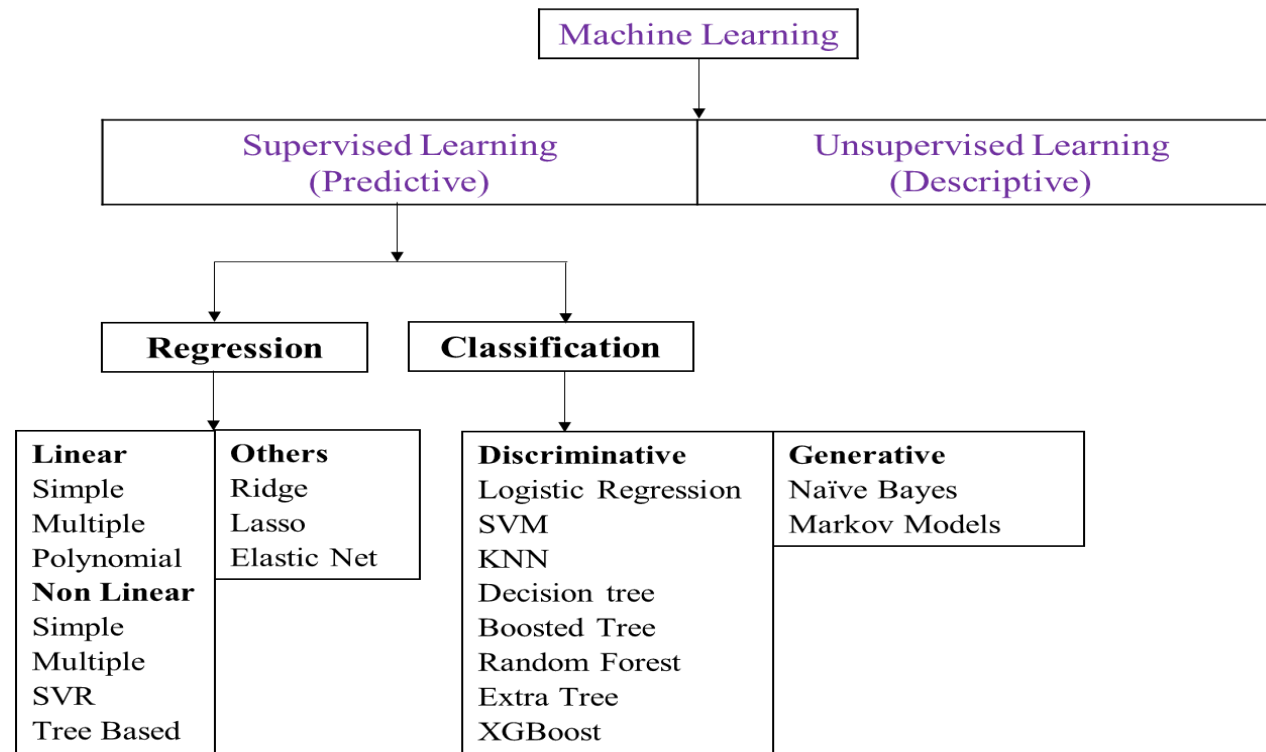
10.2.2.3 Polynomial Regression

10.2.2.4 Examples


Supervised learning is where there is an input variables (x) and an output variable (y) and an algorithm to learn the mapping function from the input to the output.

$$y = f(x)$$

The goal is to approximate the mapping function so well that when there is a new input data (x) that the model can predict the output variables (y) for that data.



Classification learning



Learning the classifier from the available data 'Training set' (Labeled)

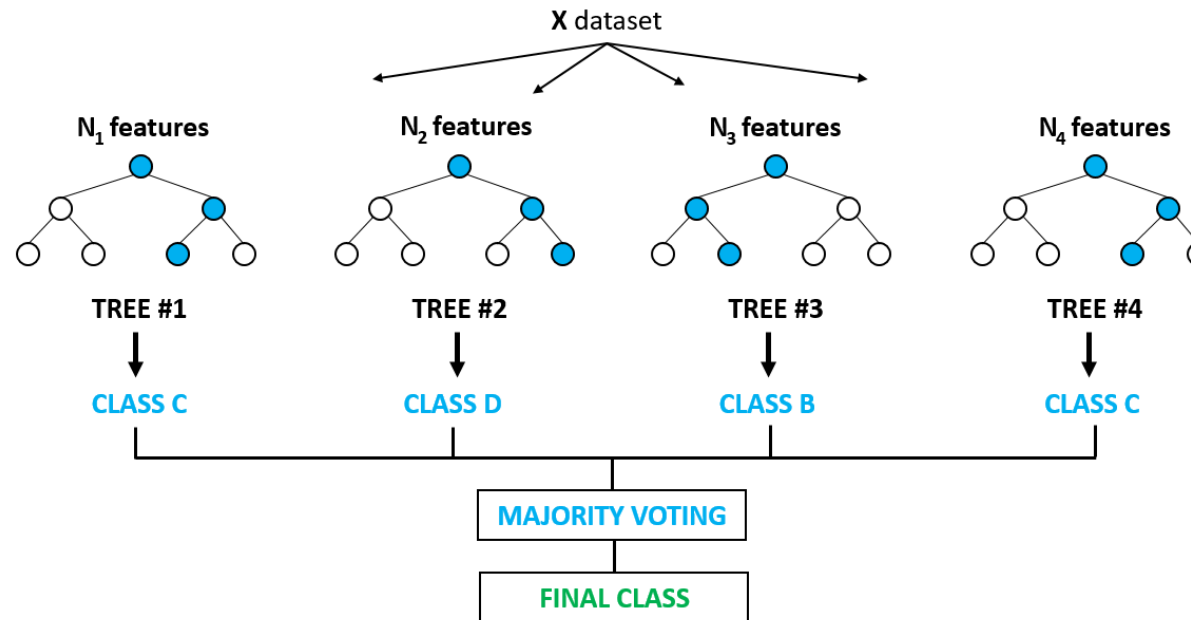
Testing how well the classifier performs 'Testing set'

Generating datasets

- Methods:
 - Holdout (2/3rd training, 1/3rd testing)
 - Cross validation (n – fold)
 - Divide into n parts
 - Train on (n-1), test on last
 - Repeat for different combinations
 - Bootstrapping
 - Select random samples to form the training set

Discriminative	Generative
Logistic Regression	Naïve Bayes
SVM	Hidden Markov Models
KNN	Gaussian Mixture Model
Decision tree	
Boosted Tree	
Random Forest	
Extra Tree	
Regularized Greedy Forest	


- An ensemble learning method for by constructing a multiple decision trees.
- The basic difference is that RF is a collection or ensemble model of numerous DT



Why Random Forest?

- One single DT would lead to over-fit model if the dataset is huge, same way like a single person might have its own perspective on the complete population.
- However if we implement the voting system and ask different individuals to interpret the data then we would be able to cover the patterns in a much meticulous way.

- A decision support tool that uses a tree like model of decisions and their possible consequences including chance of event outcomes.
- Decision tree uses divide and conquer technique for the basic learning strategy.





- Decision tree of types –
 - **Classification tree:** when the predicted outcome is the class to which the data belong.
 - **Regression tree:** when the predicted outcome is a continuous real number.
- Both the classification and regression trees have similarities as well as differences, such as procedure used to determine where to split.

Note: Classification and Regression Tree (CART) is an umbrella term used to refer both types.

Selecting Best Split

- Consider 64 examples, 29^+ and 35^-
- Which one is better?



Methods

- **Information Gain/Entropy Approach**
 - ID3 (Iterative Dichotomiser 3)
 - C4.5
 - CART (Classification and Regression Tree)
 - CHAID (CHi-squared Automatic Interaction Detector)
 - MARS
- **Gini Index Approach (*For Noisy, High Dimensional)**

Selecting Best Split

Entropy Based Approach

- The Shannon Entropy, $H(X)$ of a discrete variable X with possible values X_1, X_2, \dots, X_n and probability mass function $P(X)$ is defined as

$$H(X) = - \sum_{i=1}^n P(X_i) \ln P(X_i)$$

For Completely homogenous dataset: Entropy is 0
For Completely heterogonous dataset: Entropy is 1

- Entropy with respect to a given predictor/feature

$$E(T, X) = \sum_{i=1}^n P(X_i) H(X_i)$$

A branch with Entropy more than 1 needs splitting.
Leaf nodes have entropy 0

- Information Gain

$$IG = H(X) - E(T, X)$$

Feature/ Attribute with highest IG be the Root node.
Root node has maximum information gain.

Selecting Best Split

Gini Index Approach

- If a dataset D contains examples from n classes, Gini index of D defined as

$$gini(D) = 1 - \sum_{j=1}^n p_j^2$$

where, p_j is the relative frequency of class j in D

- If a dataset D split on A into two subset D_1 and D_2 , Gini index of D defined as

$$gini_A(D) = \frac{|D_1|}{|D|} gini(D_1) + \frac{|D_2|}{|D|} gini(D_2)$$

- Reduction in Impurity $\Delta gini(A) = gini(D) - gini_A(D)$

Handling Overfitting

- A tree that classifies the training data perfectly may not lead to best generalization performance.
 - There may be noise in the training data the tree is fitting.
 - The algorithm might be making decisions based on very little data.

A hypothesis h is said to overfit the training data


- if there is another hypothesis, h' , such that h has smaller error than h' on the training data but h has larger error on the test data than h' .

Avoiding Overfitting


- Two basic approaches
 - Prepruning: Stop growing the tree at some point when there is not enough data.
 - Postpruning: Grow the full tree and then remove nodes not having sufficient evidence.

Advantages

- What if some examples are missing values of attribute A ?
- Use training example anyway, sort through tree
 - If node n tests A , assign most common value of A among other examples sorted to node n
 - Assign most common value of A among other examples with same target value
 - Assign probability p_i to each possible value v_i of A
 - Assign fraction p_i of example to each descendant in tree
- Classify test instances with missing values in same fashion.



Classification Tree





Regression Tree

10.21.3 Logistic Regression


- Used in Classification (binary or multiclass) problem where y is a discrete value.
- Set threshold based on which classifier classifies.

Size of Tumor	Malignant
2	N
1	N
1	N
8	Y
...	...
...	...



Hypothesis

- Regression hypothesis can give values large than 1 or less than 0
- Logistic regression generates a value always between 0 or 1 (probabilistic)
- In linear regression we did $h_{\theta}(x) = (\theta^T x)$ $h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$
- For classification hypothesis representation we do $h_{\theta}(x) = g((\theta^T x))$



- If we combine these equations we can write out the hypothesis as $g(z) = 1/(1 + e^{-z})$


Where the **sigmoid** or **logistic function**

- When the probability of $y > 0.5$ then we can predict $y = 1$ Else we predict $y = 0$

$$g(z) \geq 0.5 \quad \text{when } z (= \theta^T x) \geq 0$$

$$g(z) < 0.5 \quad \text{when } z (= \theta^T x) < 0$$

- Multivariate logistic Regression $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$
- Polynomial logistic Regression $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_3 x_1^2 + \theta_4 x_2^2)$



Cost Function

- Linear regression uses the following function to determine θ

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Instead of writing the squared error term, we can write

- If we define "cost()" as $\text{cost}(h_{\theta}(x^{(i)}), y) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$

- We can **redefine J(θ)** as $J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$

Optimization

Cost Function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

Note: $y = 0$ or 1 always

Minimize cost function

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$


Gradient Descent

Repeat {


$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

} (simultaneously update all θ_j)


Multiclass Classification



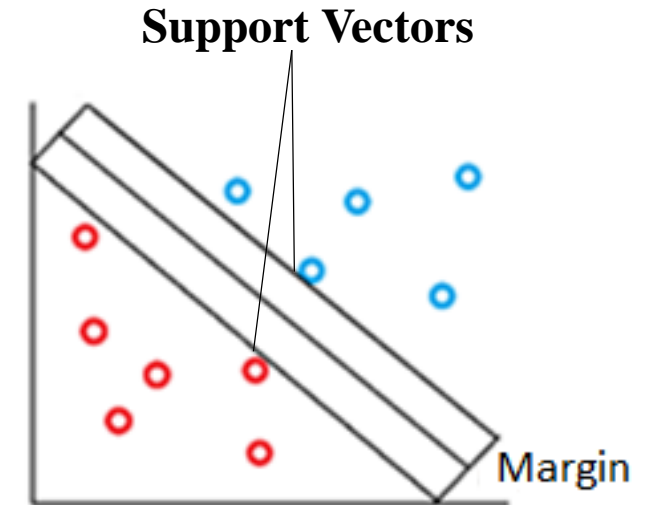
- Split the training set into three separate binary classification problems.
 - Triangle (1) vs crosses and squares (0) $h_{\theta}^1(x)$ $p(y=1 | x_1; \theta)$
 - Crosses (1) vs triangle and square (0) $h_{\theta}^2(x)$ $p(y=1 | x_2; \theta)$
 - Square (1) vs crosses and square (0) $h_{\theta}^3(x)$ $p(y=1 | x_3; \theta)$



Supervised data



Decision Boundary




Optimal Decision Boundary
(Large Margin)

Decision Boundary

- Let $\{x_1, \dots, x_n\}$ be our data set and let $y_i \in \{1, -1\}$ be the class label of x_i

- For $y_i=1$ $w^T x_i + b \geq 1$
- For $y_i=-1$ $w^T x_i + b \leq -1$

$$y_i \cdot (w^T x_i + b) \geq 1, \forall (x_i, y_i)$$



Margin

- The decision boundary should be as far away from the data of both classes as possible.
- Goal is to maximize the margin (m)

$$m = \frac{2}{\|\mathbf{w}\|}$$

Optimization

- The decision boundary should classify all points correctly for $y_i=1$

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \quad \forall i$$

- The decision boundary can be found by solving the following constrained optimization problem

$$\begin{aligned} &\text{Minimize } \frac{1}{2} \|\mathbf{w}\|^2 \\ &\text{subject to } y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \quad \forall i \end{aligned}$$

- This is a constrained optimization problem.
- Solving it requires to use Lagrange multipliers.

Optimization

- The optimization problem

$$\text{Minimize } \frac{1}{2} \|\mathbf{w}\|^2$$

$$\text{subject to } 1 - y_i(\mathbf{w}^T \mathbf{x}_i + b) \leq 0 \quad \text{for } i = 1, \dots, n$$

- The Lagrangian is

$$\mathcal{L} = \frac{1}{2} \mathbf{w}^T \mathbf{w} + \sum_{i=1}^n \alpha_i (1 - y_i(\mathbf{w}^T \mathbf{x}_i + b)) \quad \text{Note: } \|\mathbf{w}\|^2 = \mathbf{w}^T \mathbf{w}$$

Note: Lagrange multipliers are used in multivariable calculus to find maxima and minima of a function subject to constraints (like "find the highest elevation along the given path" or "minimize the cost of materials for a box enclosing a given volume").

Optimization

Gradient with respect to w and b

- Setting the gradient of L w.r.t. w and b to zero, we have

$$L = \frac{1}{2} w^T w + \sum_{i=1}^n \alpha_i (1 - y_i (w^T x_i + b))$$

$$= \frac{1}{2} \sum_{k=1}^m w^k w^k + \sum_{i=1}^n \alpha_i \left(1 - y_i \left(\sum_{k=1}^m w^k x_i^k + b \right) \right)$$

$$w + \sum_{i=1}^n \alpha_i (-y_i) x_i = 0 \Rightarrow w = \sum_{i=1}^n \alpha_i y_i x_i$$

$$\sum_{i=1}^n \alpha_i y_i = 0 \quad \begin{cases} \frac{\partial L}{\partial w^k} = 0, \forall k \\ \frac{\partial L}{\partial b} = 0 \end{cases}$$

- If we substitute $w = \sum_{i=1}^n \alpha_i y_i x_i$ to Lagrange Multipliers, we have

$$\mathcal{L} = \frac{1}{2} \sum_{i=1}^n \alpha_i y_i x_i^T \sum_{j=1}^n \alpha_j y_j x_j + \sum_{i=1}^n \alpha_i \left(1 - y_i \left(\sum_{j=1}^n \alpha_j y_j x_j^T x_i + b \right) \right)$$

$$= \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j + \sum_{i=1}^n \alpha_i - \sum_{i=1}^n \alpha_i y_i \sum_{j=1}^n \alpha_j y_j x_j^T x_i - b \sum_{i=1}^n \alpha_i y_i$$

$$= -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j + \sum_{i=1}^n \alpha_i$$

Therefore, this is a function of α_i only Since $\sum_{i=1}^n \alpha_i y_i = 0$

Optimization

Dual Problem

- The new objective function is in terms of α_i only.
- It is known as the dual problem: if we know w , we know all α_i if we know all α_i we know w .
- The original problem is known as the primal problem.

- The objective function of the dual problem needs to be maximized.

- The dual problem is therefore:
$$\begin{aligned} \max. \quad W(\alpha) &= \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1, j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j \\ \text{subject to } \alpha_i &\geq 0, \quad \sum_{i=1}^n \alpha_i y_i = 0 \end{aligned}$$

- This is a **quadratic programming (QP) problem**

- A global maximum of α_i can always be found

- \mathbf{w} can be recovered by
$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$


Optimization

Quadratic Programming Problem

- Many approaches have been proposed.
- Most are “interior-point” methods.
 - Start with an initial solution that can violate the constraints
 - Improve this solution by optimizing the objective function
- For SVM, sequential minimal optimization (SMO) seems to be the most popular.
 - A QP with two variables is trivial to solve
 - Each iteration of SMO picks a pair of (α_i, α_j) and solve the QP with these two variables; repeat until convergence
- In practice, we can just regard the QP solver as a “black-box” without bothering how it works.


Characteristics of the Solution

- Many of the α_i are zero
- So \mathbf{w} is a linear combination of a small number of data points
- \mathbf{x}_i with non-zero α_i are called support vectors (SV)
 - The decision boundary is determined only by the SV
 - Let t_j ($j=1, \dots, s$) be the indices of the s support vectors
 - We can write $\mathbf{w} = \sum_{j=1}^s \alpha_{t_j} y_{t_j} \mathbf{x}_{t_j}$
- For testing with a new data \mathbf{z}
 - Compute $\mathbf{w}^T \mathbf{z} + b = \sum_{j=1}^s \alpha_{t_j} y_{t_j} (\mathbf{x}_{t_j}^T \mathbf{z}) + b$
 - Classify \mathbf{z} as class 1 if the sum is positive else class 2



Non Linearly Separable Problems


- We allow “error” ξ_i in classification; it is based on the output of the discriminant function $\mathbf{w}^T \mathbf{x} + b$
- ξ_i approximates the number of misclassified samples




Non Linearly Separable Problems

Soft vs. Hard Margin SVM

$$\frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$$



Soft Margin SVM



Hard Margin SVM

- The algorithm try to keep ξ null, maximizing the margin.
- The algorithm does not minimize the number of error.
- Instead, it minimizes the sum of distances from the hyperplane.
- When C increases the number of errors tend to lower.
- At limit $C \rightarrow \text{infinite}$, the solution tend to that given by hard margin formulation, with 0 errors.

Non Linearly Separable Problems

Soft Margin Hyperplane



- The new conditions become

$$\begin{cases} \mathbf{w}^T \mathbf{x}_i + b \geq 1 - \xi_i & y_i = 1 \\ \mathbf{w}^T \mathbf{x}_i + b \leq -1 + \xi_i & y_i = -1 \\ \xi_i \geq 0 & \forall i \end{cases}$$


- ξ_i are “slack variables” in optimization
 - Note that $\xi_i=0$ if there is no error for \mathbf{x}_i
 - ξ_i is an upper bound of the number of errors
-
- We want to minimize $\frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$
subject to $y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0$
-
- C : tradeoff parameter between error and margin

Transformation


- Generalizing large-margin classifier with non linear decision
- Idea is to transform \mathbf{x}_i to a higher dimensional space
 - Input space: the space the point \mathbf{x}_i are located
 - Feature space: the space of $f(\mathbf{x}_i)$ after transformation
- Why transform?
 - Classification can become easier with a proper transformation.
 - In the XOR problem, for ex. adding a new feature of $x_1 x_2$ make the problem linearly separable.



Transformation



- Feature space is of higher dimension than the input space in practice
- Computation in the feature space can be costly because it is high dimensional.
- The kernel trick comes to rescue.



Input space

Feature space

(Hyperplane in R^2 is a line) (Hyperplane in R^3 is a plane)

The Kernel Trick

$$\begin{aligned} \max. \quad W(\alpha) &= \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1, j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j \\ \text{subject to } C &\geq \alpha_i \geq 0, \sum_{i=1}^n \alpha_i y_i = 0 \end{aligned}$$

- Recall the SVM optimization problem
- The data points only appear as **inner product**
- As long as we can calculate the inner product in feature space, we don't need mapping explicitly.
- Many common geometric operations (angles, distances) can be expressed by inner products.

- Suppose $f(\cdot)$ is given as follows $\phi\left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}\right) = (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, x_2^2, \sqrt{2}x_1x_2)$
- An inner product in the feature space is $\langle \phi\left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}\right), \phi\left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix}\right) \rangle = (1 + x_1y_1 + x_2y_2)^2$
- So, if we define the kernel function as follows, there is no need to carry out $f(\cdot)$ explicitly


$$K(\mathbf{x}, \mathbf{y}) = (1 + x_1y_1 + x_2y_2)^2$$

- This use of kernel function to avoid carrying out $f(\cdot)$ explicitly is known as the **kernel trick**.

- Define the kernel function K by $K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$

The Kernel Trick

Kernel Function



- In practical use of SVM, user specifies kernel function; transformation $f(\cdot)$ is not explicitly stated.
- Given a kernel function $K(\mathbf{x}_i, \mathbf{x}_j)$, the transformation $f(\cdot)$ is given by its Eigen functions.
 - Eigen functions can be difficult to construct explicitly.
 - This is why people only specify kernel function without worrying about exact transformation.
- Another view: kernel function, being an inner product, is a similarity measure between the objects.




SVM Parameter Tuning

- **Kernel:** It specifies the kernel type to be used in the algorithm.
 - It can be 'linear', 'poly', 'rbf', 'sigmoid'
 - The default value is 'rbf'
- **Degree:** The degree of the polynomial kernel function ('poly')
- **C:** It is the regularization parameter of the error term.
- **Gamma:** It is the kernel coefficient for 'rbf', 'poly', and 'sigmoid'.

<https://cs.stanford.edu/people/karpathy/svmjs/demo/>

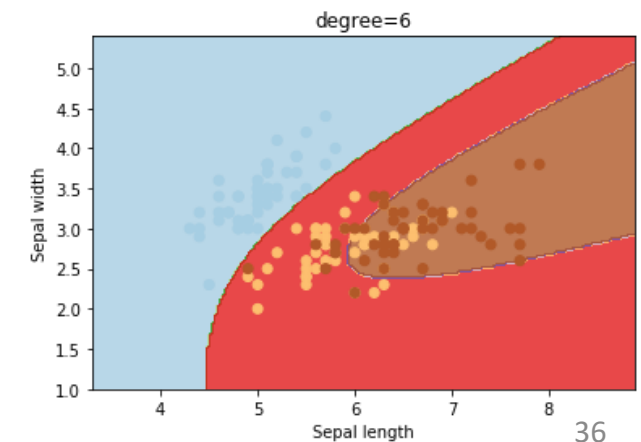
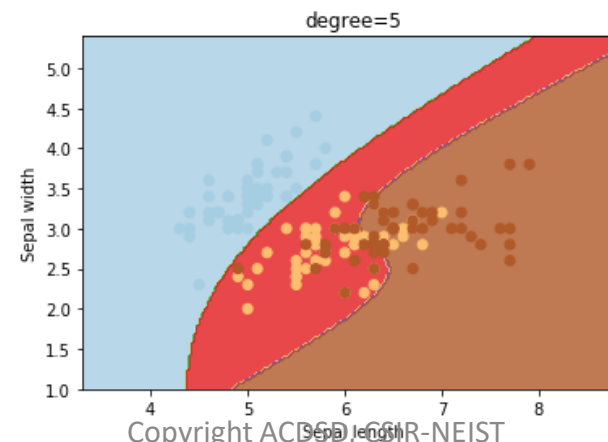




SVM Parameter Tuning

- kernel parameters selects the type of hyperplane used to separate the data.
- ‘linear’ use a linear hyperplane.
- ‘rbf’ and ‘poly’ uses a non linear hyper-plane.





SVM Parameter Tuning

- Degree is a parameter used when kernel is set to 'poly'.
- It's basically the degree of the polynomial used to find the hyperplane to split the data.



SVM Parameter Tuning

- C Value controls the trade off between smooth boundary and classifying training points correctly.
- Increasing C values may lead to getting more training points correctly.
- With that overfitting the training data.



- Finding out smooth decision boundary vs correct (with overfitting) is part of artistry of ML.
- So try different values of c to get the perfectly balanced curve and avoid over fitting.

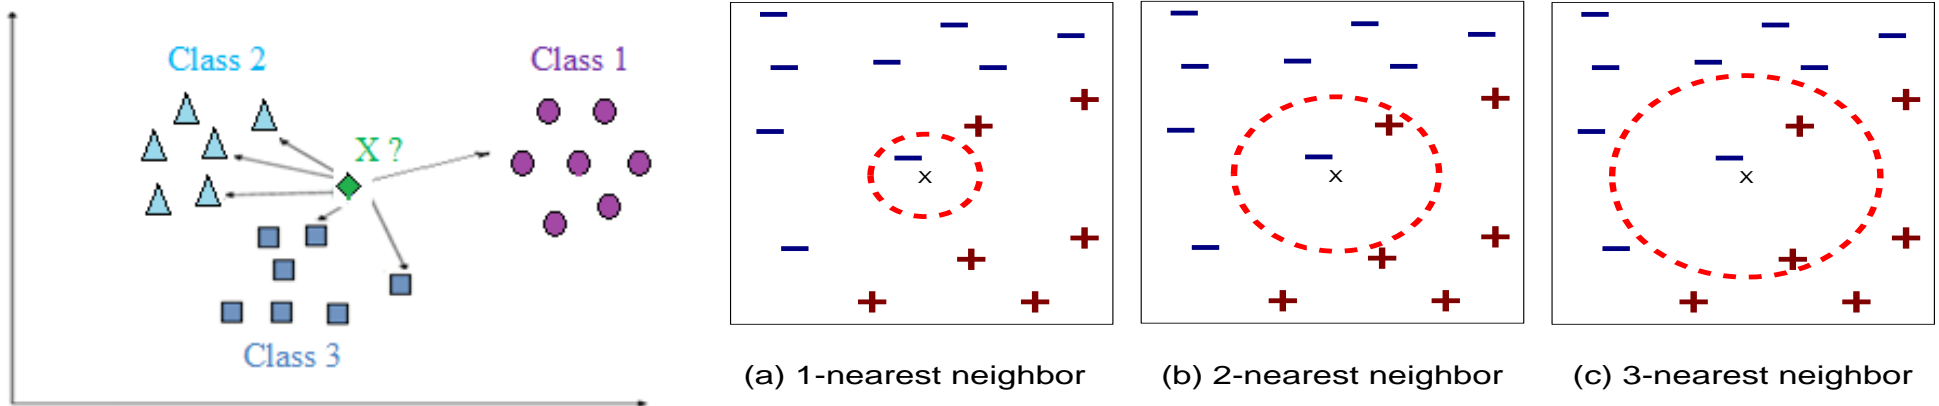
SVM Parameter Tuning

- Gamma defines how far data points are considered in calculation of separation line.



- For low gamma even the far away points get considerable weight and it results more linear curve.
- For high gamma the closer points get more weight and it results in a wiggly curve.

- Explains a categorical value using the majority votes of (k) nearest neighbors.
- Select the tuning parameter k based on data and experiment.



Algorithm

1. Load the data and initialize the value of k
2. For getting the predicted class, iterate from 1 to total number of training data points
 - a) Calculate distance (metrics ex.- Euclidean) between test data and each row of training data.
 - b) Sort the calculated distances in ascending order based on distance values
 - c) Get top k rows from the sorted array
 - d) Get the most frequent class of these rows
 - e) Return the predicted class

Choosing K


Choosing the value of k :

- If k is too small, sensitive to noise points
- If k is too large, neighborhood may include points from other classes

Rule of thumb:

$K = \text{sqrt}(N)$ where N : number of training points

In practice, k is usually chosen to be odd, so as to avoid ties.



Laziness


- **Training:** $O(1)$
- **Testing:** $O(nd)$

Training set: n sample

Testing instance: 1 sample

Nearest Neighbor: Compare One by one to each training instance

Complexity: N comparisons and each comparisons takes d operations



Making Fast

- **Reduce d :** dimensionality reduction
- **Reduce n :** do not compare to all training examples

Bayes' Theorem

The diagram shows the equation $P(c | x) = \frac{P(x | c)P(c)}{P(x)}$ with four labels and arrows pointing to its parts: 'Likelihood' points to $P(x | c)$, 'Class Prior Probability' points to $P(c)$, 'Posterior Probability' points to $P(c | x)$, and 'Predictor Prior Probability' points to $P(x)$.

$$P(c | x) = \frac{P(x | c)P(c)}{P(x)}$$

Labels and arrows:

- Likelihood (points to $P(x | c)$)
- Class Prior Probability (points to $P(c)$)
- Posterior Probability (points to $P(c | x)$)
- Predictor Prior Probability (points to $P(x)$)

$$P(c | X) = P(x_1 | c) \times P(x_2 | c) \times \cdots \times P(x_n | c) \times P(c)$$

- Probabilistic Classifier based on Baye's theorem with strong (Naïve) independence assumptions among the features.

Why Naïve?

- It assumes the features that go into the model is independent of each other.
- That is changing the value of one feature, does not directly influence any other features.
- Because of independence nature it makes assumptions that may or may not turn out to be correct.

Multivariant Bayesian Classifier

- Compute the posterior probability $P(C | A_1, A_2, \dots, A_n)$ for all values of C using the Bayes theorem

$$P(C | A_1, A_2, \dots, A_n) = \frac{P(A_1, A_2, \dots, A_n | C) P(C)}{P(A_1, A_2, \dots, A_n)}$$

- Choose value of C that maximizes $P(C | A_1, A_2, \dots, A_n)$
- Equivalent to choosing value of C that maximizes $P(A_1, A_2, \dots, A_n | C) P(C)$

Note: Other names "Simple Baye's" and "Independence Baye's"

Variants

Based on the probability distribution adopted

Bernoulli: Assumes that all features are binary such that they take only two values.

Multinomial: Assumes that features element represents frequency i.e. the number of times it appears.

Gaussian: Based on a continuous distribution and it's suitable for more generic classification tasks.

Logistic Regression.

Example: A patient is admitted with abdominal sepsis (blood poisoning). The case is severe enough to warrant surgery. The patient is wheeled to the operation theatre. Let us speculate what will happen after the surgery. Let

$$\begin{aligned} Y &= 1 \text{ if death follows surgery,} \\ &= 0 \text{ if the patient survives} \end{aligned}$$

The variables X_1 , X_2 , X_3 , and X_5 are categorical covariates and X_4 is a continuous covariate. The categorical variables are binary with only two possible values. It is felt that the outcome Y depends on these covariates.

Response variable

Y

Covariates, predictors or independent variables

X_1, X_2, X_3, X_4, X_5

Basic question: Which of the predictors have an impact on the response?

```
dim(Sepsis) # dimension of the data set
[1] 106 6
```

```
head(Sepsis) # Display the head
```

	Shock	Malnutrition	Alcoholism	Age	Infarction	Death
1	0	0	0	56	0	0
2	0	0	0	80	0	0
3	0	0	0	61	0	0

```
Summary(Sepsis) # summary statistics of each variable
```

Shock	Malnutrition	Alcoholism	Age
Min. :0.00000	Min. :0.0000	Min. :0.0000	Min. :17.00
1st Qu.:0.00000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:33.00
Median :0.00000	Median :0.0000	Median :0.0000	Median :52.50
Mean :0.09434	Mean :0.3019	Mean :0.2075	Mean :51.28
3rd Qu.:0.00000	3rd Qu.:1.0000	3rd Qu.:0.0000	3rd Qu.:68.75
Max. :1.00000	Max. :1.0000	Max. :1.0000	Max. :94.00

```
apply(Sepsis, 2, class) # class of each variable
```

Shock	Malnutrition	Alcoholism	Age	Infarction	Outcome
"integer"	"integer"	"integer"	"integer"	"integer"	"integer"

Let us convert each binary variable into a factor (categorical variable with direction). Create another copy of Sepsis

```
Sepsis1 <- Sepsis
> Sepsis1$Shock <- as.factor(Sepsis1$Shock)
> Sepsis1$Malnutrition <- as.factor(Sepsis1$Malnutrition)
> Sepsis1$Alcoholism <- as.factor(Sepsis1$Alcoholism)
> Sepsis1$Infarction <- as.factor(Sepsis1$Infarction)
> Sepsis1$Death <- as.factor(Sepsis1$Death)
```

```
summary(Sepsis1) # summary statistics
```

Shock	Malnutrition	Alcoholism	Age	Infarction	Death
0:96	0:74	0:84	Min. :17.00	0:92	Min. :0.0000
1:10	1:32	1:22	1st Qu.:33.00	1:14	1st Qu.:0.0000
			Median :52.50		Median :0.0000
			Mean :51.28		Mean :0.1981
			3rd Qu.:68.75		3rd Qu.:0.0000
			Max. :94.00		Max. :1.0000

Classification Trees

Kyphosis is a misalignment of vertebrate. Surgery is helpful to align the problematic vertebrate. Different parameters such as age, number of the topmost vertebrae surgery was done, outcome of surgery was taken into consideration.

10.21.7 Examples

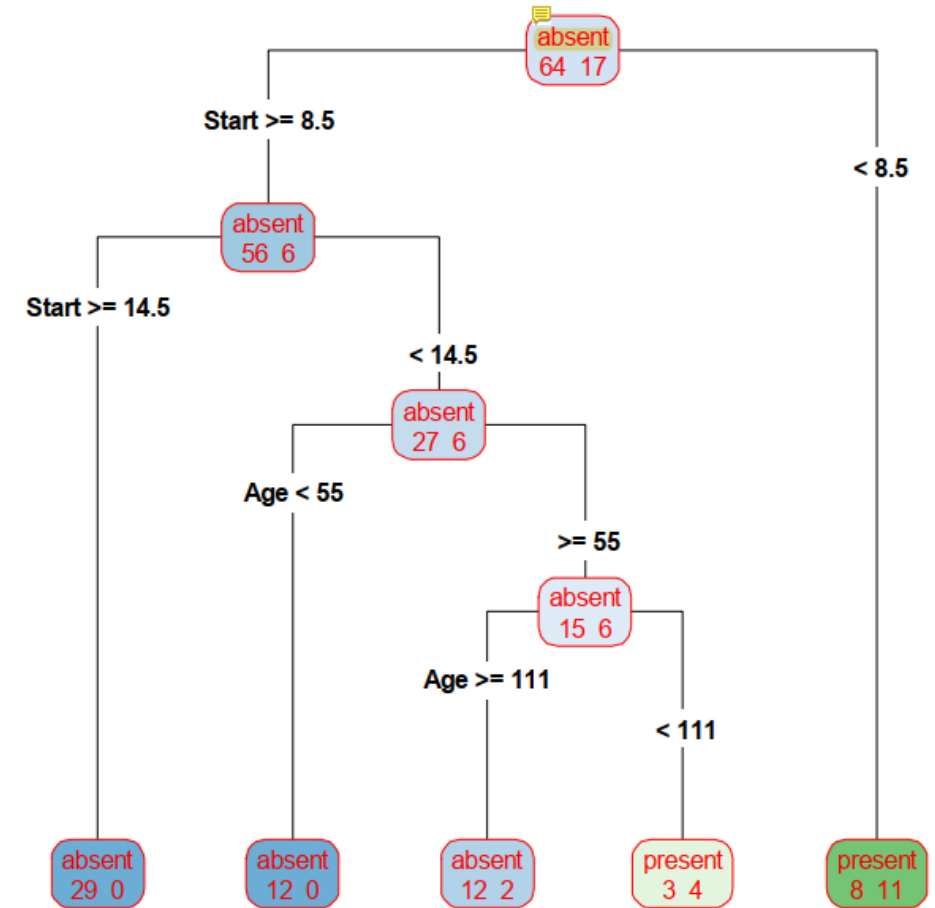
```
library(rpart) #Activate the rpart library contains different dataset
library("rpart.plot") #Activate the plot function
data(kyphosis) # access kyphosis data set
dim(kyphosis)
[1] 81 4
head(kyphosis)
```

```
> head(kyphosis)
  Kyphosis Age Number Start
1  absent  71      3      5
2  absent 158      3     14
3  present 128      4      5
4  absent   2      5      1
5  absent   1      4     15
6  absent   1      2     16
```

```
MB <- rpart(Kyphosis ~ Age+Number+Start,data=kyphosis)
rpart.plot(MB, type = 4, extra = 1, digits = 3, col = "red")
```

Description of the tree:

1. The root node has 81 subjects, for 64 of them kyphosis is absent and 17 present.
2. All those subjects with $\text{Start} \geq 8.5$ go into the left node. Total number of subjects in the left node is 62, 56 of them have kyphosis absent.
3. All those subjects with $\text{Start} < 8.5$ go into the right node. Total number of subjects in the right node is 19, 8 of them have kyphosis absent.
4. This node is a terminal node. No further split is envisaged because the total number of observations is $19 \leq 20$. The command stops splitting a node if the size of the node is 20 or less (default). This is a pruning strategy. This node is declared 'present' as per the 'majority rule' paradigm.



Misclassification rate = re-substitution error
 $= 100 * (8 + 0 + 2 + 3) / 81 = 16\%$

RandomForest

Objective: The response variable is ‘Species.’ The flower could be any one of the three types. I have a flower. I do not know to which species the flower belongs. I know its petal length, petal width, sepal length, and sepal width. Guess the species. This is a classification problem or a pattern recognition problem.

```
>data(iris)      # Download the data.  
> head(iris)     # Understand the data. The top six rows of the data ...
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

```
➤dim(iris)
```

```
[1] 150 5
```

```
> table(iris$Species)
```

setosa	versicolor	virginica
50	50	50

Install and activate the package ‘randomForest.’

```
install.packages(“randomForest”)
```

```
Library(randomForest)
```

```
MB <- randomForest(Species ~ ., data = iris, importance = T)
```

```
print(MB)    # print the output
```

Call:

```
randomForest(formula = Species ~ ., data = iris, importance = T,
proximity = T)
```

Type of random forest: classification

Number of trees: 500

No. of variables tried at each split: 2

(Here the number of covariates is 4 and m = 2.)

OOB estimate of error rate: 4%

Confusion matrix:

	setosa	versicolor	virginica	
class.error				
setosa	50	0	0	0.00
versicolor	0	47	3	0.06
virginica	0	3	47	0.06


One can obtain a graph of the importance of the predictors in the classification problem.

```
> varImpPlot(MB, pch = 16, col = "red", n.var = 4, sort = T, main = "Importance of Variables for the Iris data")
```

Understand the output:

1. R says that this is a classification problem.
2. R produces a forest of 500 trees.
3. There are four predictors. At every node, it chose 2 predictors at random for splitting the node.
4. The given data were poured into the root node for classification. The acronym 'OOB' stands for 'out-of-bag.' The OOB (error) estimate is 4%.


The confusion matrix spells out where errors of misclassification occurred



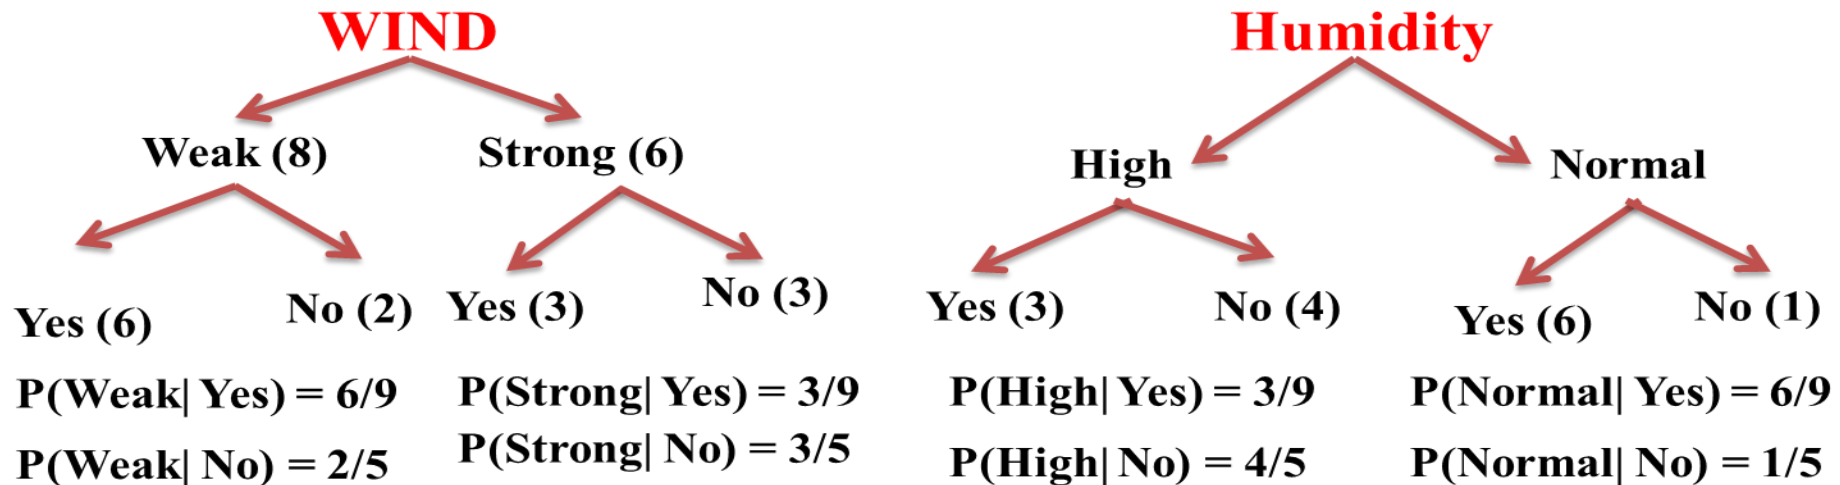
10.21.7 Examples

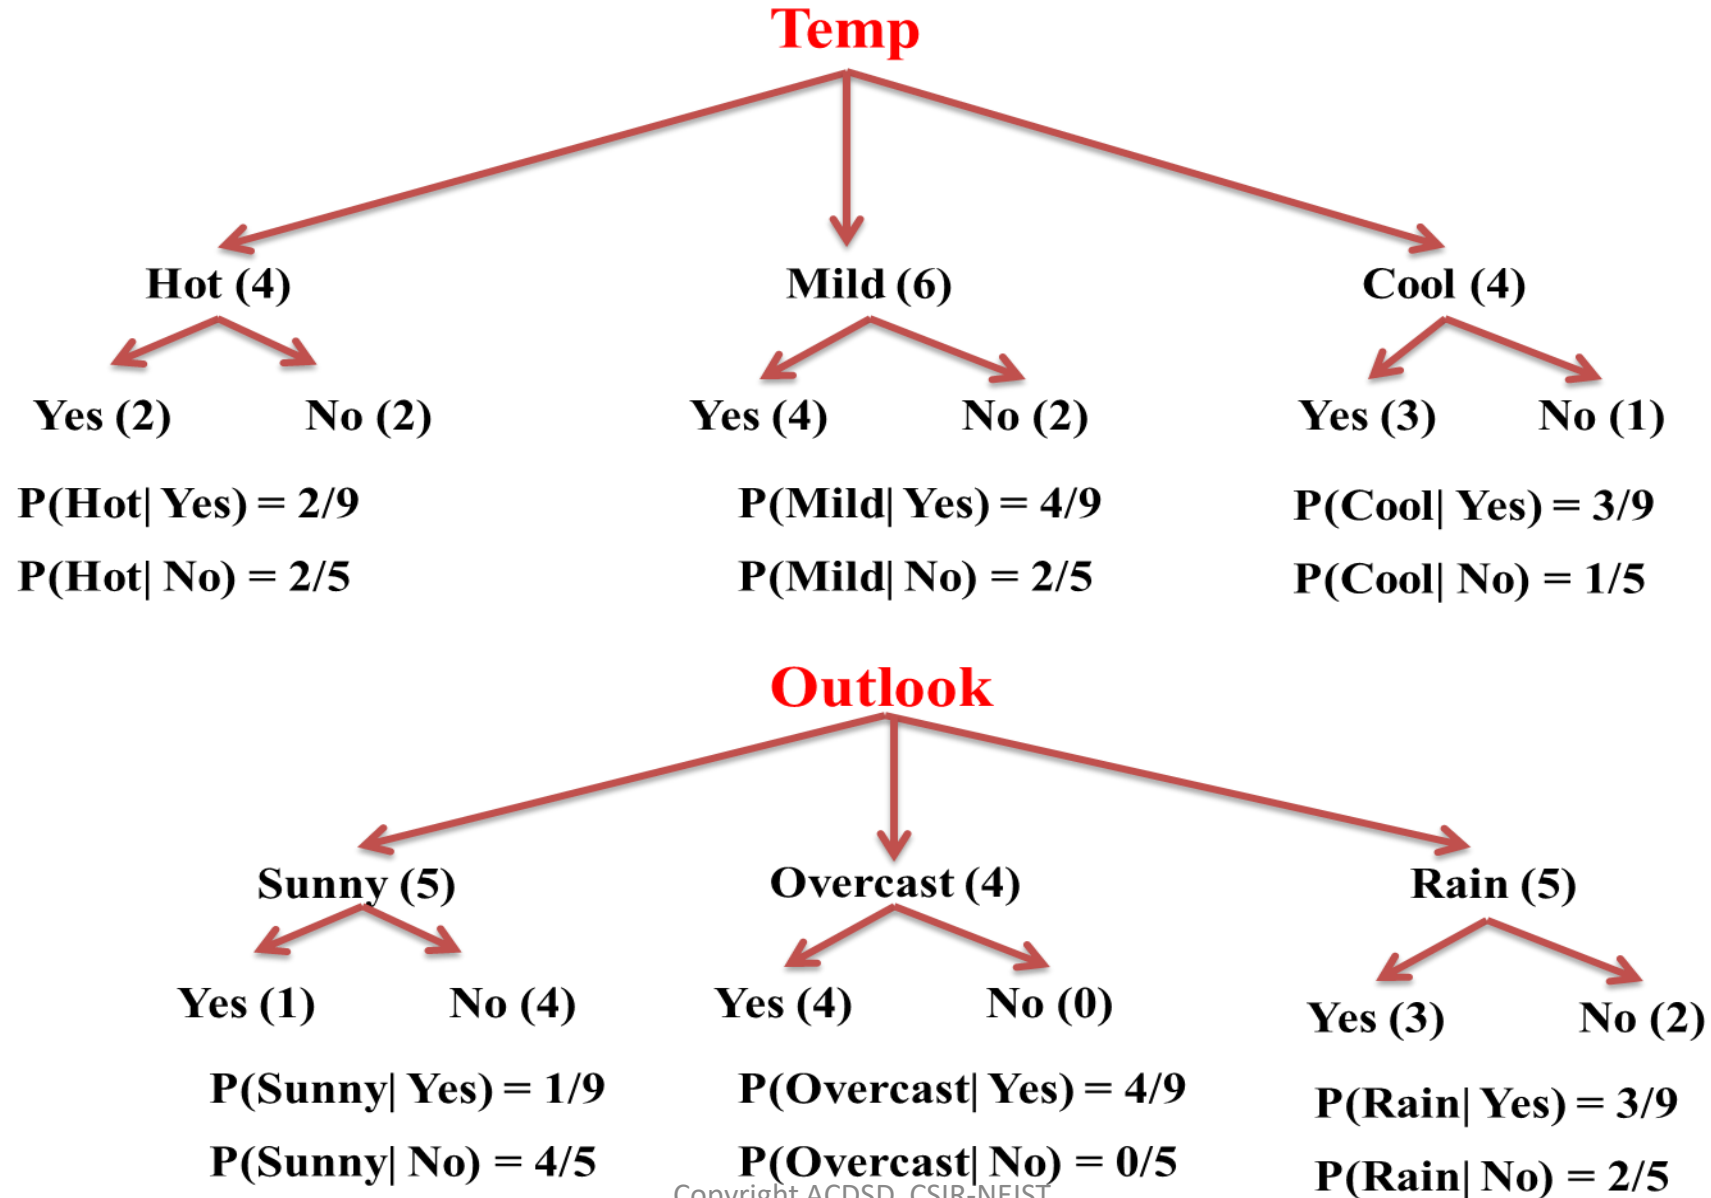
**Find the probability to play on 15th day where conditions are
temp = cool, humidity = high wind = strong and outlook = sunny**

Day	Outlook	Temp	Humidity	Wind	Play
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
3	Overcast	Hot	High	Weak	Yes
4	Rain	Mild	High	Weak	Yes
5	Rain	Cool	Normal	Weak	Yes
6	Rain	Cool	Normal	Strong	No
7	Overcast	Cool	Normal	Strong	Yes
8	Sunny	Mild	High	Weak	No
9	Sunny	Cool	Normal	Weak	Yes
10	Rain	Mild	Normal	Weak	Yes
11	Sunny	Mild	Normal	Strong	Yes
12	Overcast	Mild	High	Strong	Yes
13	Overcast	Hot	Normal	Weak	Yes
14	Rain	Mild	High	Strong	No



Expansion from right to left





Temp	Outlook	Wind	Humidity
$P(\text{Hot} \text{Yes}) = 2/9$	$P(\text{Sunny} \text{Yes}) = 1/9$	$P(\text{Weak} \text{Yes}) = 6/9$	$P(\text{High} \text{Yes}) = 3/9$
$P(\text{Hot} \text{No}) = 2/5$	$P(\text{Sunny} \text{No}) = 4/5$	$P(\text{Weak} \text{No}) = 2/5$	$P(\text{High} \text{No}) = 4/5$
$P(\text{Mild} \text{Yes}) = 4/9$	$P(\text{Overcast} \text{Yes}) = 4/9$	$P(\text{Strong} \text{Yes}) = 3/9$	$P(\text{Normal} \text{Yes}) = 6/9$
$P(\text{Mild} \text{No}) = 2/5$	$P(\text{Overcast} \text{No}) = 0/5$	$P(\text{Strong} \text{No}) = 3/5$	$P(\text{Normal} \text{No}) = 1/5$
$P(\text{Cool} \text{Yes}) = 3/9$	$P(\text{Rain} \text{Yes}) = 3/9$		
$P(\text{Cool} \text{No}) = 1/5$	$P(\text{Rain} \text{No}) = 2/5$		

Let $X = \{\text{Sunny, Cool, High, Strong}\}$

$$\begin{aligned}
 P(X|\text{Yes}) &= P(\text{Yes}) * P(\text{Sunny}|\text{Yes}) * P(\text{Cool}|\text{Yes}) * P(\text{High}|\text{Yes}) * P(\text{Strong}|\text{Yes}) \\
 &= 9/14 * 1/9 * 3/9 * 3/9 * 3/9 \\
 &= 0.0053
 \end{aligned}$$

$$\begin{aligned}
 P(X|\text{No}) &= P(\text{No}) * P(\text{Sunny}|\text{No}) * P(\text{Cool}|\text{No}) * P(\text{High}|\text{No}) * P(\text{Strong}|\text{No}) \\
 &= 5/14 * 4/5 * 1/5 * 4/5 * 3/5 \\
 &= 0.0206
 \end{aligned}$$

$$P(X|\text{No}) > P(X|\text{Yes})$$

Classification with Naïve Bayes

A dataset containing 400 observations with 4 variables GRE, GPA, rank of school where a candidate has studied and admission status. Using this dataset design a classification model with Naïve Bayes classifier in R.

```
'data.frame':  400 obs. of  4 variables:
 $ admit: int  0 1 1 1 0 1 1 0 1 0 ...
 $ gre  : int  380 660 800 640 520 760 560 400 540 700 ...
 $ gpa   : num  3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
 $ rank  : int  3 3 1 4 4 2 1 2 3 2 ...
```

#libraries

```
library(naivebayes)
library(dplyr)
library(ggplot2)
library(psych)
```

data

```
data <- read.csv(file.choose(), header = T)
str(data)
xtabs(~admit+rank, data = data)
data$rank <- as.factor(data$rank)
data$admit <- as.factor(data$admit)
```

→ Load R libraries

→ Convert the data into factor

#visualization

```
pairs.panels(data[-1])
data %>%
  ggplot(aes(x=admit, y=gpa, fill = admit)) +
  geom_boxplot() +
  ggtitle("Box Plot")
```

Box plot on the
gpa data



```
data %>% ggplot(aes(x=gpa, fill = admit)) +
  geom_density(alpha=0.8, color= 'black') +
  ggtitle("Density Plot")
```

Density plot
on the gpa data

Data Partition

```
set.seed(1234)
ind <- sample(2, nrow(data), replace = T, prob = c(0.8,
0.2))
train <- data[ind == 1,]
test <- data[ind == 2,]
```

Partitioning data
for training and
test



Naive Bayes Model

```
model <- naive_bayes(admit ~ ., data = train, ,
usekernel = T)
```

```
model
```

```
train %>%
```

```
  filter(admit == "0") %>%
```

```
  summarise(mean(gre), sd(gre))
```

Naïve Bayes
model for training

A priori probabilities:

```
      0      1
0.6861538 0.3138462
```

Tables:

```
::: gre::0 (KDE)
```

```
::: gre::1 (KDE)
```

Call:

```
density.default(x = x, na.rm = TRUE)
```

Data: x (223 obs.); Bandwidth 'bw' = 35.5

```
      x      y
Min.   :193.5 Min.   :6.010e-07
1st Qu.:371.7 1st Qu.:2.924e-04
Median :550.0 Median :1.291e-03
Mean   :550.0 Mean   :1.401e-03
3rd Qu.:728.3 3rd Qu.:2.405e-03
Max.   :906.5 Max.   :3.199e-03
```

Call:

```
density.default(x = x, na.rm = TRUE)
```

Data: x (102 obs.); Bandwidth 'bw' = 39.59

```
      x      y
Min.   :181.2 Min.   :1.145e-06
1st Qu.:365.6 1st Qu.:2.007e-04
Median :550.0 Median :1.129e-03
Mean   :550.0 Mean   :1.354e-03
3rd Qu.:734.4 3rd Qu.:2.751e-03
Max.   :918.8 Max.   :3.465e-03
```

```
::: gpa::0 (KDE)
```

Call:

```
density.default(x = x, na.rm = TRUE)
```

Data: x (223 obs.); Bandwidth 'bw' = 0.1134

```
      x      y
Min.   :2.080 Min.   :0.0002229
1st Qu.:2.645 1st Qu.:0.0924939
Median :3.210 Median :0.4521795
Mean   :3.210 Mean   :0.4419689
3rd Qu.:3.775 3rd Qu.:0.6603271
Max.   :4.340 Max.   :1.1433285
```

```
::: gpa::1 (KDE)
```

Call:

```
density.default(x = x, na.rm = TRUE)
```

Data: x (102 obs.); Bandwidth 'bw' = 0.1234

```
      x      y
Min.   :2.25 Min.   :0.0005231
1st Qu.:2.78 1st Qu.:0.0800747
Median :3.31 Median :0.4801891
Mean   :3.31 Mean   :0.4710851
3rd Qu.:3.84 3rd Qu.:0.8626207
Max.   :4.37 Max.   :1.0595464
```

```
::: rank (Categorical)
```

```
rank      0      1
1 0.10313901 0.24509804
2 0.36771300 0.42156863
3 0.33183857 0.24509804
4 0.19730942 0.08823529
```

Prediction

```
p <- predict(model, train, type = 'prob')
head(cbind(p, train))
```

Confusion Matrix - train data

```
p1 <- predict(model, train)
(tab1 <- table(p1, train$admit))
1 - sum(diag(tab1)) / sum(tab1)
```

Confusion Matrix - test data

```
p2 <- predict(model, test)
(tab2 <- table(p2, test$admit))
1 - sum(diag(tab2)) / sum(tab2)
```

→ Prediction

→ Confusion matrix
of train and test
data

	0	1	admit	gre	gpa	rank
1	0.8528794	0.1471206	0	380	3.61	3
2	0.5621460	0.4378540	1	660	3.67	3
3	0.2233490	0.7766510	1	800	4.00	1
4	0.8643901	0.1356099	1	640	3.19	4
6	0.6263274	0.3736726	1	760	3.00	2
7	0.5933791	0.4066209	1	560	2.98	1

Confusion matrix of train data

```
p1      0      1
0 203    69
1   20    33
> 1 - sum(diag(tab1)) / sum(tab1)
[1] 0.2738462
```

Confusion matrix of test data

```
p2      0      1
0  47    20
1   3     5
> 1 - sum(diag(tab2)) / sum(tab2)
[1] 0.3066667
```

Regression can of many types based on –

- Parameters
 - Linear Regression (parameters are linear)
 - Non-linear (parameters are non linear)
- Number of independent Variables
 - Simple
 - Multiple
 - Polynomial

Linear	Non Linear	Others
Simple	Simple	SVR
Multiple	Multiple	Regression Tree
Polynomial (special case)		

- Regression is a statistical technique used to establish a relationship model between two variables.
- Originally invented by Francis Galton to study the relationships between parents and children.
- One of these variable is called predictor variable whose value is gathered through experiments.
- The other variable is called response variable whose value is derived from the predictor variable.

Dependent/Measured/Outcome/Experimental/Response - what we are trying to predict
Independent/Covariate/Feature/Explanatory/Predictor - what we are using to predict

- Gather a sample (input) of response and corresponding predictor variable by experimentation.
- Find out a relationship between the response variables and corresponding output variable.
- Based on the fitted relationship/line (model) we predict the unknown predictor.

Regression can of many types –


- Linear Regression (parameters are linear)
- Non-linear (parameters are non linear)
- Simple
- Multiple
- Polynomial (Curvy linear – a special case of multiple regression)

- Training set

X	Y
2.1	4
1.4	2.5
1.7	3.5
.85	1
...	...
...	...

- Notation

- m = number of **training examples**
- x 's = input variables / features
- y 's = output variable "target" variables
- (x, y) - single training example
- (x^i, y^j) - specific example (i^{th} training example)



$$m = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$
$$c = \bar{y} - m\bar{x}$$

Hypothesis

- Take training set
- Pass into a learning algorithm
 - Algorithm outputs a function (denoted h) (h = **hypothesis**)

$$h_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 \mathbf{x}$$

Y is a linear function of x
 θ_i are **parameters** - θ_0 is zero condition or slope θ_1 is gradient

- This function (hypothesis) takes an input X (e.g. size of new house)
 - Uses parameters determined by a learning system
 - Tries to output the estimated value of Y based on that input
- This kind of linear regression with one variable is called **univariate linear regression**.

Error Evaluation Metrics

$$MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}|$$

\hat{y} – predicted value of y
 \bar{y} – mean value of y

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

$$RMSE = \sqrt{MSE} = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2}$$

$$R^2 = 1 - \frac{\sum (y_i - \hat{y})^2}{\sum (y_i - \bar{y})^2}$$

$$R_{adj}^2 = 1 - (1 - R^2) \frac{N-1}{N-M-1}$$

M – no of features

Cost Function

- Choosing values for θ_i (parameters) - Different values give you different hypothesis.
- Based on training set generate parameters which give best hypothesis (h)

- **The squared error:** $\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

- **Cost function:** $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

1/m means we determine the average
The 2 makes the math a bit easier

Helps out how to fit the best parameters/hypothesis by minimizing squared error.


Optimization

Gradient Descent Algorithm

- GD Algorithm to minimize cost function used all over machine learning for minimization.
- Minimizing J means minimizing θ_0/θ_1 i.e. we get the values of θ_0 and θ_1 which find on average the minimal deviation of x from y

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Problem
 - We have $J(\theta_0, \theta_1)$
 - We want to get **$\min J(\theta_0, \theta_1)$**
- Gradient descent applies to more general functions
 - $J(\theta_0, \theta_1, \theta_2 \dots \theta_n)$
 - $\min J(\theta_0, \theta_1, \theta_2 \dots \theta_n)$




- Linear models are frequently favorable due to their interpretability and often good predictive performance.
- Ordinary Least Squares (OLS) estimation faces challenges: *Overfitting* and *Interpretability*
- Fitting techniques alternative to OLS:
 - Subset selection
 - Dimension reduction
 - Shrinkage methods
- Fit linear models with least squares but impose constraints on the coefficients
- Instead, alternative formulations add a penalty in the OLS formula
- Best known are ridge regression and LASSO (least absolute shrinkage operator)
 - Ridge regression can shrink parameters close to zero
 - LASSO models can shrink some parameters exactly to zero
 - Performs implicit variable selection

$$\text{OLS Estimation: } \boldsymbol{\beta}_{\text{OLS}} = \min_{\boldsymbol{\beta}} \text{RSS} = \min_{\boldsymbol{\beta}} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$

$$\text{Ridge Regression: } \boldsymbol{\beta}_{\text{ridge}} = \min_{\boldsymbol{\beta}} \underbrace{\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2}_{\text{RSS}} + \underbrace{\lambda \sum_{j=1}^p \beta_j^2}_{\text{shrinkage penalty}}$$

$$\text{Least Absolute Shrinkage and Selection Operator (LASSO): } \boldsymbol{\beta}_{\text{LASSO}} = \min_{\boldsymbol{\beta}} \underbrace{\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2}_{\text{RSS}} + \underbrace{\lambda \sum_{j=1}^p |\beta_j|}_{\text{shrinkage penalty}}$$



- How do we fit the model to this data
- To map our old linear hypothesis and cost functions to these polynomial descriptions set

$$x_1 = x$$

$$x_2 = x^2$$

$$x_3 = x^3$$

The Topics Covered in This Section

10.3.1 Clustering

10.3.1.1 K-Means

10.3.1.2 K Nearest Neighbors

10.3.1.3 Association Rule Learning

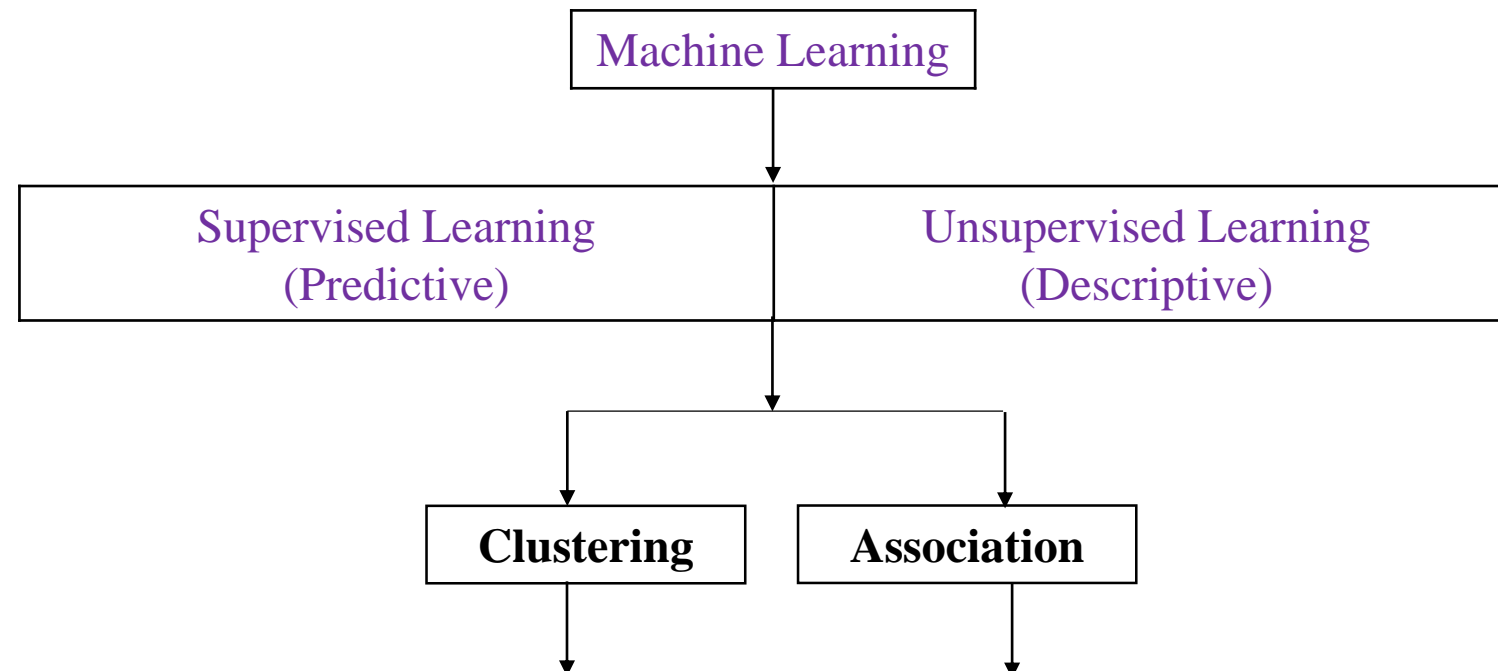
10.3.2 Dimensionality Reduction

10.3.2.1 PCA



10.3.2.2 LDA

10.3.2.3 GDA

10.3.2.4 Examples



Partitional K-means K-modes K-median K-medoids PAM CLARA CLARANS Fuzzy C-means	Hierarchy AGNES DIANA BIRCH CURE ROCK Chameleon Hierarchical k-means	Density DBSCAN OPTICS Mean-shift DENCLUE Graph CLICK MST	Grid STING CLIQUE WaveCluster Dstream OptiGrid	Model COBWEB GMM SOM ART Autoencoder Distribution DBCLASD GMM	Apriori Eclat FPGrowth
---	---	---	--	--	------------------------------



$\{X1, X2\} \rightarrow \{Y1, Y2\}$
 Antecedent Consequent


Itemset = $\{X1, X2, Y1, Y2\}$

- Clustering involves grouping data into categories based on some measure of similarity or distance.
- Thus clustering requires some methods for computing (dis)**similarity** or **distance** between each pair of observations.
- The result of this computation is known as a dissimilarity or **distance metrics**.

Similarity among points is defined using –

- Some inter-observation distance measures
- Correlation-based distance measures

- Clustering algorithm for partitioning a given data set into a set of k groups (i.e. k clusters), where k represents the number of groups pre-specified by the analyst.
- K-means algorithm allows that each one of the data belong to only a cluster.
- Therefore, this algorithm is a definite clustering and hard clustering algorithm.



Divides the observations into discrete groups based on some distance metric

ALGORITHM: K-Means

Input: $\{x^1, x^2, x^3 \dots, x^n\}$ //Training Set

K // Number of desired clusters

Output: K // Set of Clusters

Randomly initialize K cluster centroids as $\{\mu_1, \mu_2, \mu_3 \dots \mu_K\}$

Repeat

for $i=1$ to n

$c^{(i)}$ = index of cluster centroid (from 1 to K) closest to $x^{(i)}$

for $k=1$ to K

μ_k = average(mean) of points assigned to cluster k

Until convergence criteria is meet

- Optimization objective: $J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m ||x^{(i)} - \mu_{c^{(i)}}||^2$

where,

- μ_k is the cluster associated with centroid k
- c^i is the index of clusters $\{1, 2, \dots, K\}$ to which x^i is currently assigned.
- So, μ_{c^i} is the cluster centroid of the cluster to which example x^i has been assigned to.

- Minimizing cost:
$$\min_{\substack{c^{(1)}, \dots, c^{(m)}, \\ \mu_1, \dots, \mu_K}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

where,

- The **cluster assigned step** is minimizing $J(\dots)$ with respect to $c^1, c^2 \dots c^i$
- The **move centroid step** is minimizing $J(\dots)$ with respect to μ

Initializing the centroids – Random Initialization

- Randomly pick K training examples
- Set μ_1 up to μ_K to these example's values

Effects of Random Initialization


- Kmeans can converge to different solutions depending on the initialization setup
- That means risk of local optimum.
- The local optimum are valid convergence, but local optimum are not global ones
- If this is a concern go for multiple random initializations.

Deciding K - Elbow method

- Compute cost function at a range of K values (As K increases $J(\dots)$ minimum value decreases)
- Plot (K vs $J()$)
- Look for the "elbow" on the graph and Chose the "elbow" number of clusters.

Risks:

- Normally we don't get a nice line – no clear elbow on curve
- Not really that helpful




Strength:

- Easy to implement.
- Efficient: $O(tkn)$ where n is number of objects, k is number of clusters, and t is number of iterations.
- k-Means produce tighter clusters than hierarchical clustering.
- K-means is scalable and efficient for large dataset with complexity $O(nkt)$.


Weakness:

- Need to specify k , the number of clusters, in advance.
- Initial seeds have a strong impact on the final results.
- Applicable only when mean is defined - Categorical data.
- Unable to handle noisy data and outliers.
- Not suitable to discover clusters with non-convex shapes.


- Input features – can be both quantitative and qualitative
- Output feature – always categorical values
- Explains a categorical value using the majority votes of (k) nearest neighbors.
- Select the tuning parameter k based on data and experiment.
- Requires distance metric to define proximity between any two data points –
 - Manhattan Distance
 - Euclidean Distance
 - Minkowski Distance
 - Hamming Distance
 - Cosine Distance
 - Mahalanobis Distance



(a) 1-nearest neighbor



(b) 2-nearest neighbor




(c) 3-nearest neighbor

Distance Metrics

Chebyshev Distance:	$\max_i (x_i - y_i)$	Canberra:	$\sum_{i=1}^m \frac{ x_i - y_i }{ x_i + y_i }$
Manhattan/City Block Distance (L_∞):	$\sum_{i=1}^k x_i - y_i $	Correlation:	$\frac{\sum_{i=1}^m (x_i - \bar{x}_i)(y_i - \bar{y}_i)}{\sqrt{\sum_{i=1}^m (x_i - \bar{x}_i)^2 \sum_{i=1}^m (y_i - \bar{y}_i)^2}}$
Euclidean Distance:	$\sqrt{\sum_{i=1}^k (x_i - y_i)^2}$	Chi:	$\sum_{i=1}^m \frac{1}{sum_i} \left(\frac{x_i}{size_x} - \frac{y_i}{size_y} \right)^2$
Minkowski Distance:	$\left(\sum_{i=1}^k (x_i - y_i)^q \right)^{1/q}$		
Hamming Distance:	$D_H = \sum_{i=1}^k x_i - y_i $		
Cosine Distance:	$similarity = \cos(\theta) = \frac{A \cdot B}{\ A\ \ B\ } = \frac{\sum_{i=1}^n A_i \times B_i}{\sqrt{\sum_{i=1}^n (A_i)^2} \times \sqrt{\sum_{i=1}^n (B_i)^2}}$		
Mahalanobis Distance:	$d(\vec{x}, \vec{y}) = \sqrt{(\vec{x} - \vec{y})^T S^{-1} (\vec{x} - \vec{y})}$	where S^{-1} is covariance matrix of the input data	

KNN Algorithm

1. Load the data
2. Initialize the value of k
3. For getting the predicted class, iterate from 1 to total number of training data points
 - a) Calculate distance (metrics ex.- Euclidean) between test data and each row of training data.
 - b) Sort the calculated distances in ascending order based on distance values
 - c) Get top k rows from the sorted array
 - d) Get the most frequent class of these rows
 - e) Return the predicted class



What is k=?


Choosing the value of k:

- If k is too small, sensitive to noise points
- If k is too large, neighborhood may include points from other classes

Rule of thumb:

$K = \sqrt{N}$ where N: number of training point

In practice, k is usually chosen to be odd, so as to avoid ties.



- **Training:** $O(1)$

- **Testing:** $O(nd)$

Training set: n sample

Testing instance: 1 sample

Nearest Neighbor:

Compare One by one to each training instance

Making Fast

Complexity: N comparisons and each comparisons takes d operations

- **Reduce d:** dimensionality reduction
- **Reduce n:** do not compare to all training examples

What is Association Mining: Finding/mining frequent patterns and rules.

Stages:

- Find frequent patterns.
- Derive associations from frequent patterns.

Example (Applications):

- Graphs - Social Network Analysis
- Transactions - Market Basket Analysis
- Sequences - Time Series Analysis
- Topic Identification - Co-occurrence of words, Plagiarism Detection
- Bio-maskers - Genes or proteins vs. disease

Association Rule

- An implication expression of the form $X \rightarrow Y$, where X and Y are itemsets
- Example: $\{\text{Milk, Diaper}\} \rightarrow \{\text{Beer}\}$

Rule Evaluation Metrics:

- Support (s): Fraction of transactions that contain both X and Y
- Confidence (c): Measures how often items in Y appear in transactions that contain X

$$\text{support} = \frac{(X \cup Y).count}{n}$$

$$\text{confidence} = \frac{(X \cup Y).count}{X.count}$$

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

$\{\text{Milk, Diaper}\} \Rightarrow \text{Beer}$

$$s = \frac{\sigma(\text{Milk, Diaper, Beer})}{|T|} = \frac{2}{5} = 0.4$$

$$c = \frac{\sigma(\text{Milk, Diaper, Beer})}{\sigma(\text{Milk, Diaper})} = \frac{2}{3} = 0.67$$

- The goal of association rule mining is to find frequent itemset i.e. all rules having
 - Support $\geq \textit{minsup}$ threshold
 - Confidence $\geq \textit{minconf}$ threshold

Method

- Brute-force approach
- Apriori
- Frequency Pattern Growth

Brute Force Approach

1. Frequent Itemset Generation
 - Generate all combinatorically possible itemsets whose support \geq minsup
2. Rule Generation
 - List all possible association rules.
 - Compute the support and confidence for each rule.
 - Prune rules that fail the *minsup* and *minconf* thresholds.

Note: Frequent itemset generation is still computationally expensive

Factors Affecting Complexity:

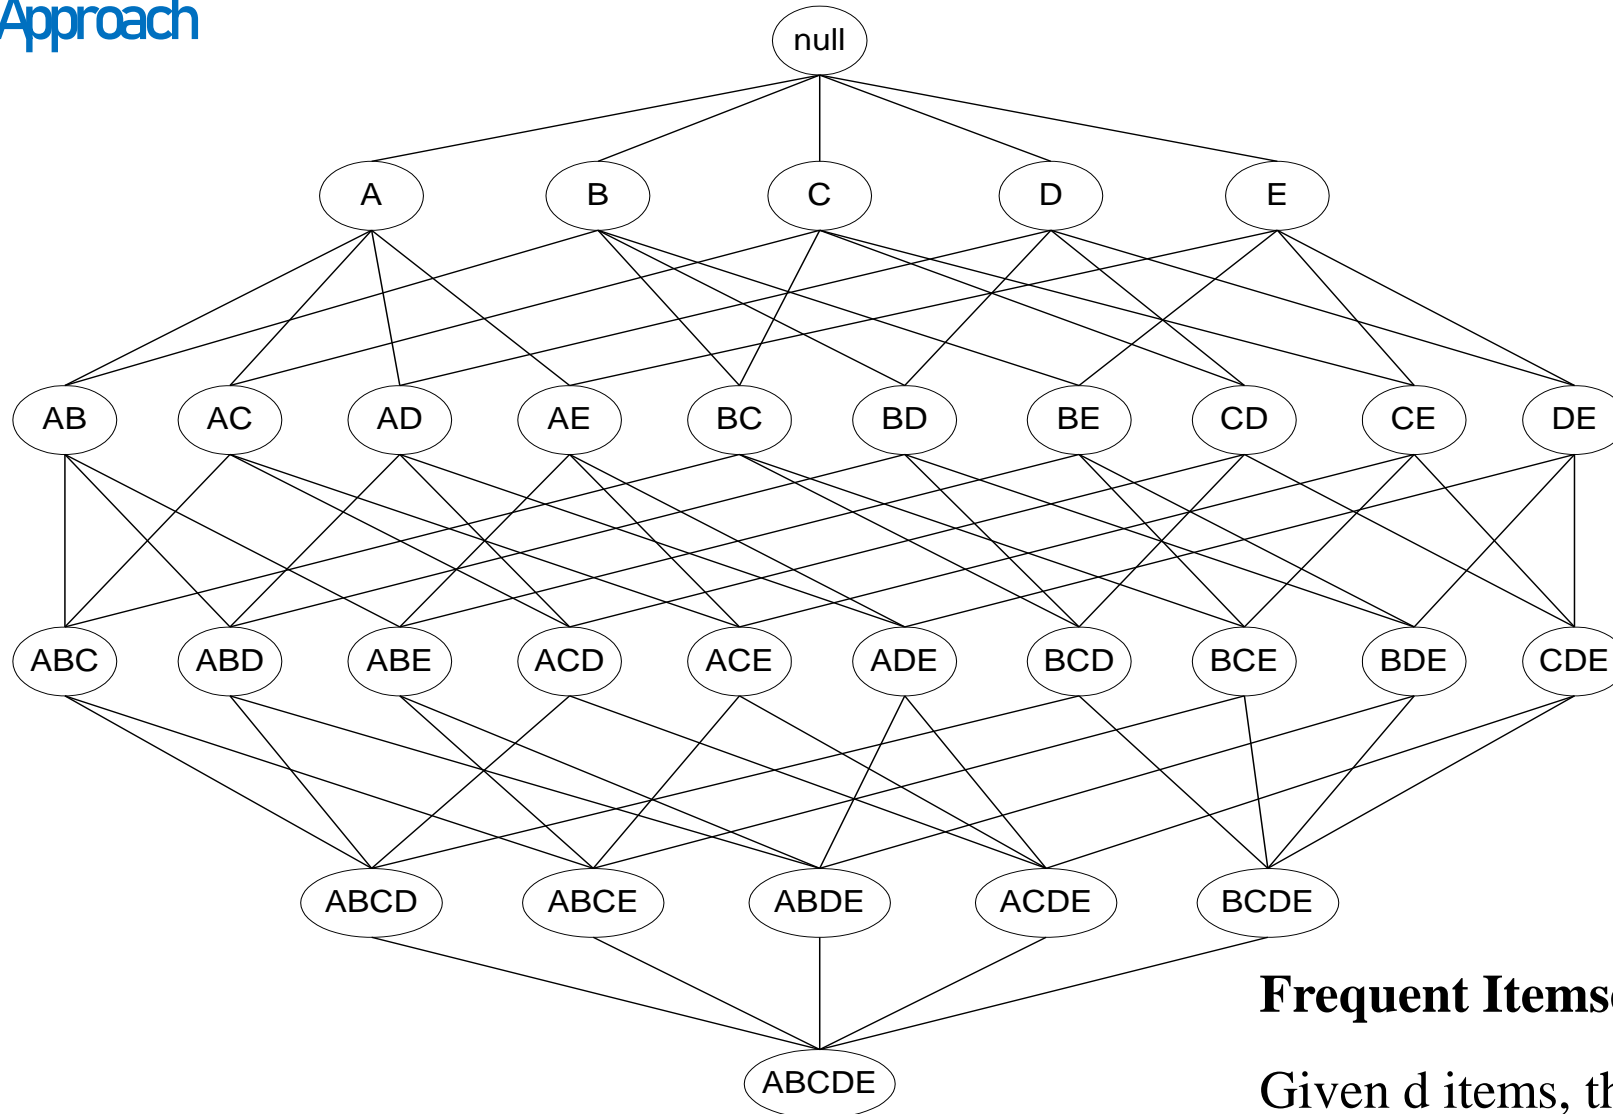
- Choice of minimum support threshold
- Dimensionality (number of items) of the data set
- Size of database (number of transaction)
- Average transaction width

Brute Force Approach

Rule Generation

- How to efficiently generate rules from frequent itemsets?
- In general, confidence does not have an anti-monotone property
 $c(ABC \rightarrow D)$ can be larger or smaller than $c(AB \rightarrow D)$
- But confidence of rules generated from the same itemset has an anti-monotone property
e.g., $L = \{A, B, C, D\}$:
$$c(ABC \rightarrow D) \geq c(AB \rightarrow CD) \geq c(A \rightarrow BCD)$$
- Confidence is anti-monotone w.r.t. number of items on the RHS of the rule

Brute Force Approach



Frequent Itemset Generation

Given d items, there are 2^d possible candidate itemsets

Apriori Algorithm

Goal

- Reducing Number of Candidates.

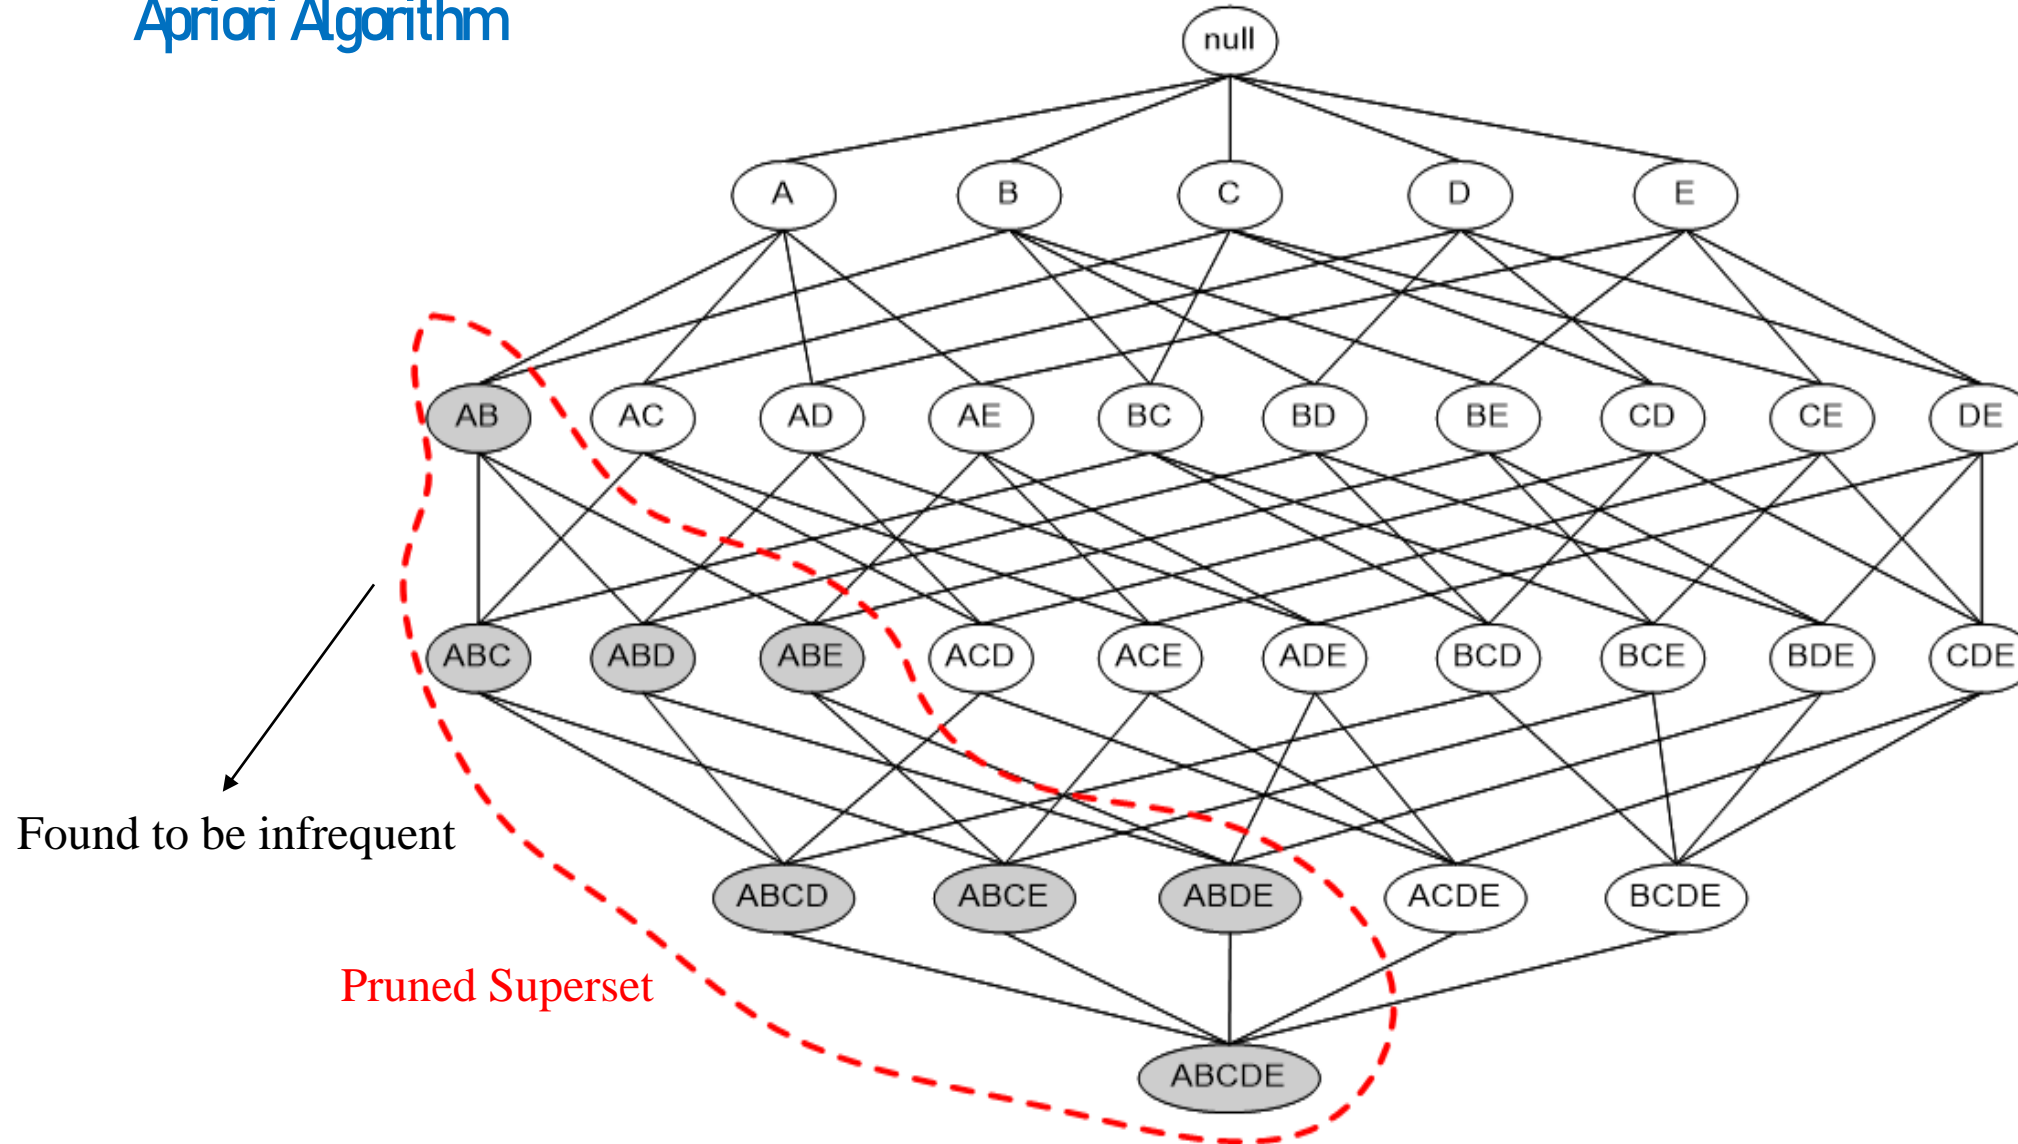
Apriori principle

- If an itemset is frequent, then all of its subsets must also be frequent.
- Apriori principle holds due to the following property of the support measure:

$$\forall X, Y : (X \subseteq Y) \Rightarrow s(X) \geq s(Y)$$

- Support of an itemset never exceeds the support of its subsets (anti-monotone property of support)

Apriori Algorithm



Apriori Algorithm

Pattern Evaluation

- Association rule algorithms tend to produce too many rules
 - Many of them are uninteresting or redundant
 - Redundant if $\{A,B,C\} \rightarrow \{D\}$ and $\{A,B\} \rightarrow \{D\}$ have same support & confidence
- Interestingness measures (Statistical-based) can be used to prune/rank the derived patterns.
- In the original formulation of association rules, support & confidence are the only measures used.

Apriori Algorithm

Statistical Independence

- Population of 1000 students
 - 600 students know how to swim (S)
 - 700 students know how to bike (B)
 - 420 students know how to swim and bike (S,B)
- $P(S \wedge B) = 420/1000 = 0.42$
- $P(S) \times P(B) = 0.6 \times 0.7 = 0.42$
- $P(S \wedge B) = P(S) \times P(B) \Rightarrow$ Statistical independence
- $P(S \wedge B) > P(S) \times P(B) \Rightarrow$ Positively correlated
- $P(S \wedge B) < P(S) \times P(B) \Rightarrow$ Negatively correlated

Apriori Algorithm

Statistical-based Measures

- Take into account statistical dependence

$$Lift = \frac{P(Y | X)}{P(Y)}$$

$$Interest = \frac{P(X, Y)}{P(X)P(Y)}$$

$$PS = P(X, Y) - P(X)P(Y)$$

$$\phi - coefficient = \frac{P(X, Y) - P(X)P(Y)}{\sqrt{P(X)[1 - P(X)]P(Y)[1 - P(Y)]}}$$

Contingency table for supports $X \rightarrow Y$

	Coffee	<u>Coffee</u>	
Tea	15	5	20
<u>Tea</u>	75	5	80
	90	10	100

Association Rule: Tea \rightarrow Coffee

Confidence = $P(\text{Coffee} | \text{Tea}) = 0.75$

but $P(\text{Coffee}) = 0.9$

\Rightarrow Lift = $0.75/0.9 = 0.8333$ (< 1 , so negatively associated)

FP Growth

- An efficient and scalable method to complete set of frequent patterns.
- It allow frequent item set discovery without candidate item set generation.
- Two step approach:
 1. Build a compact data structure called the FP-Tree.
 2. Extracts frequent item set directly from the FP-Tree.

Steps

1. Calculate minimum support
2. Find frequency of occurrence of each item
3. Prioritize the items according to the frequencies
4. Order the transactions according to priority of items
5. Build FP Tree (row by row for each ordered transaction)
6. Validation of the FP Tree
7. Extract Information

Example on Frequent Patterns

Suppose we have the following transaction

TID	Items
1	E, A, D, B
2	D, A, C, E, B
3	C, A, B, E
4	B, A, D
5	D
6	D, B
7	A, D, E
8	B, C

Step 1

Step 2

Items	Freq
A	5
B	6
C	3
D	6
E	4

Step 3

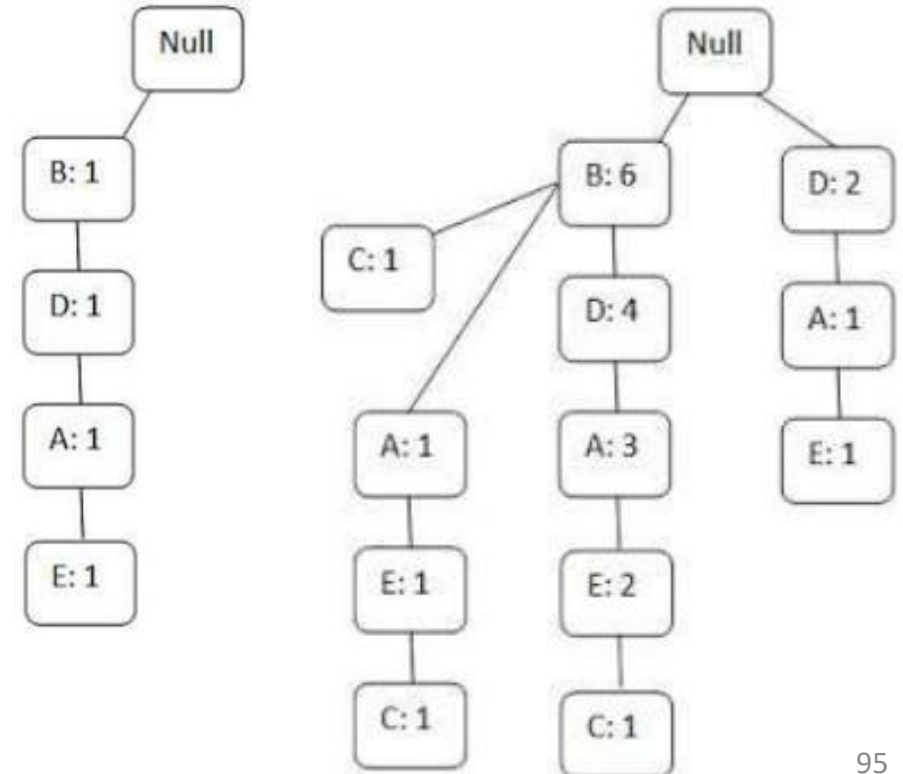
Items	Frequency	Priority
A	5	3
B	6	1
C	3	5
D	6	2
E	4	4

Example on Frequent Patterns

Step 4


TID	Items	Ordered
1	E, A, D, B	B, D, A, E
2	D, A, C, E, B	B, D, A, E, C
3	C, A, B, E	B, A, E, C
4	B, A, D	B, D, A
5	D	D
6	D, B	B, D
7	A, D, E	D, A, E
8	B, C	B, C

Step 5



Example on Frequent Patterns


Step 6



Items	Frequency
A	5
B	6
C	3
D	6
E	4

Example on Frequent Patterns

Step 7



Priority item	Item	Conditional pattern base
B (1)	C	{ B: 1 } {BAE:1} {BDAE:1}
D (2)	E	{BA:1} {BDA:2} {DA:1}
A (3)	A	{B:1} {BD:3} {D:1}
E (4)	D	{B:4}
C (5)	B	-



Priority item	Item	Conditional pattern base		Conditional FP tree (>minsupp)	Frequent pattern
B (1)	C	{ B: 1 } {BAE:1} {BDAE:1}	{B:3, A: 2, E:2, D:1}	{B:3}	{BC:3}
D (2)	E	{BA:1} {BDA:2} {DA:1}	{A:3, B:3,D:2}, {D: 1, A:1}	{A:3, B:3}	{AE: 3} {BE:3} {ABE:3}
A (3)	A	{B:1} {BD:3} {D:1}	{B:4, D:3} {D:1}	{B:4, D:3}	{BA:4} {DA: 3} (BAD:3)
E (4)	D	{B:4}	{B:4}	{B:4}	{BD:4}
C (5)	B	-	-	-	

- **Purpose**
 - Avoid curse of dimensionality.
 - Allow data to be more easily visualized.
 - May help to eliminate irrelevant (unimportant, dummy, redundant) features or reduce noise.
 - Reduce amount of time and memory required by data mining algorithms.
- **Techniques** (Reducing feature set with minimal loss of information)
 - Principal Component Analysis (PCA)
 - Linear Discriminant Analysis (LDA)
 - Generalized Discriminant Analysis (GDA)
- **Curse of Dimensionality**
 - Issue due to high dimension data
 - More data is good, more detailed (dimension) data may not be

- PCA is an unsupervised dimensionality reduction algorithm as it ignores the class labels.
- PCA does not select a set of features and discard other features.
- PCA infers some new features, which best describe the type of class from the existing features.
- PCA finds the lines of maximum variance in the dataset known as principal components.
 - PC1 having the maximum variance
 - PC2 having second maximum variance and so on.
- The components are orthogonal and hence highly uncorrelated with each other.
- PCA works on eigenvectors and eigenvalues of the covariance matrix (linear combination of variables), which is the equivalent of fitting those straight PC lines to the variance of the data.

- LDA is supervised dimension reduction algorithm as it considers ignores the class labels.
- LDA is not necessarily a classifier, but can be used as one as it maintains the class separability.
- LDA finds components (linear combination of the observed variables) that maximize class separation (between-class variance) when such prior information is available (i.e. supervised).
- LDA finds a centroid of each class datapoints.
- Based on this, it determines a new dimension which should satisfy two criteria:
 - Maximize the distance between the centroid of each class.
 - Minimize the variation within each category.
- If the variation of data is minimum, then less overlapping between the classes will be there and maximum separation will be maintained between the different classes.

PCA vs LDA



Note: PC values are between -1 and +1

- Generalized Discriminant Analysis (GDA) has provided an extremely powerful approach to extracting non-linear features
- The Generalized Discriminant Analysis is used for multi-class classification problems.
- GDA is employed to reduce the high dimensional data vectors and identification is handled in a low dimensional space with high efficiency and low use of system resources
- ***Virtues where LDA fails:***
 - Linear boundaries fail to separate data
 - May want to model non linear data
 - Too many parameters estimated with high variance
 - Need of regularization
- Generalization of LDA to solve the issues:
 - ❑ Flexible Discriminant Analysis (FDA): LDA in enlarged space of predictors
 - ❑ Penalized Discriminant Analysis (PDA) : Can be used with large number of predictors
 - ❑ Mixture Discriminant Analysis (MDA): Model by mixture of two or more Gaussians sharing same covariance matrix

Cluster Analysis with Hierarchical & K-Means

A dataset containing 22 observations with 9 variables for different companies has been given. Using this dataset design clustering models with Hierarchical & K-Means clustering approaches in R.

```
'data.frame': 22 obs. of 9 variables:
 $ Company      : chr  "Arizona " "Boston " "Central " "Commonwealth" ...
 $ Fixed_charge : num  1.06 0.89 1.43 1.02 1.49 1.32 1.22 1.1 1.34 1.12 ...
 $ RoR          : num  9.2 10.3 15.4 11.2 8.8 13.5 12.2 9.2 13 12.4 ...
 $ Cost         : int  151 202 113 168 192 111 175 245 168 197 ...
 $ Load        : num  54.4 57.9 53 56 51.2 60 67.6 57 60.4 53 ...
 $ D.Demand     : num  1.6 2.2 3.4 0.3 1 -2.2 2.2 3.3 7.2 2.7 ...
 $ Sales        : int  9077 5088 9212 6423 3300 11127 7642 13082 8406 6455 ...
 $ Nuclear      : num  0 25.3 0 34.3 15.6 22.5 0 0 0 39.2 ...
 $ Fuel_Cost    : num  0.628 1.555 1.058 0.7 2.044 ...
```

reading data from dataset

```
mydata <- read.csv(file.choose(), header=T)
str(mydata)
head(mydata)
tail(mydata)
```

Read first and last 6 entries

```
> head(mydata) #first 6 data
  Company Fixed_charge RoR Cost Load D.Demand Sales Nuclear Fuel_Cost
1  Arizona      1.06  9.2  151 54.4      1.6  9077      0.0    0.628
2  Boston       0.89 10.3  202 57.9      2.2  5088     25.3    1.555
3  Central      1.43 15.4  113 53.0      3.4  9212      0.0    1.058
4 Commonwealth  1.02 11.2  168 56.0      0.3  6423     34.3    0.700
5  Con Ed NY    1.49  8.8  192 51.2      1.0  3300     15.6    2.044
6  Florida      1.32 13.5  111 60.0     -2.2 11127     22.5    1.241

> tail(mydata) #last 6 data
  Company Fixed_charge RoR Cost Load D.Demand Sales Nuclear Fuel_Cost
17 San Diego      0.76  6.4  136 61.9      9.0  5714      8.3    1.920
18 Southern      1.05 12.6  150 56.7      2.7 10140      0.0    1.108
19 Texas         1.16 11.7  104 54.0     -2.1 13507      0.0    0.636
20 Wisconsin    1.20 11.8  148 59.9      3.5  7287     41.1    0.702
21 United States  1.04  8.6  204 61.0      3.5  6650      0.0    2.116
22 Virginia      1.07  9.3  174 54.3      5.9 10093     26.6    1.306
```

10.3.2.4 Examples


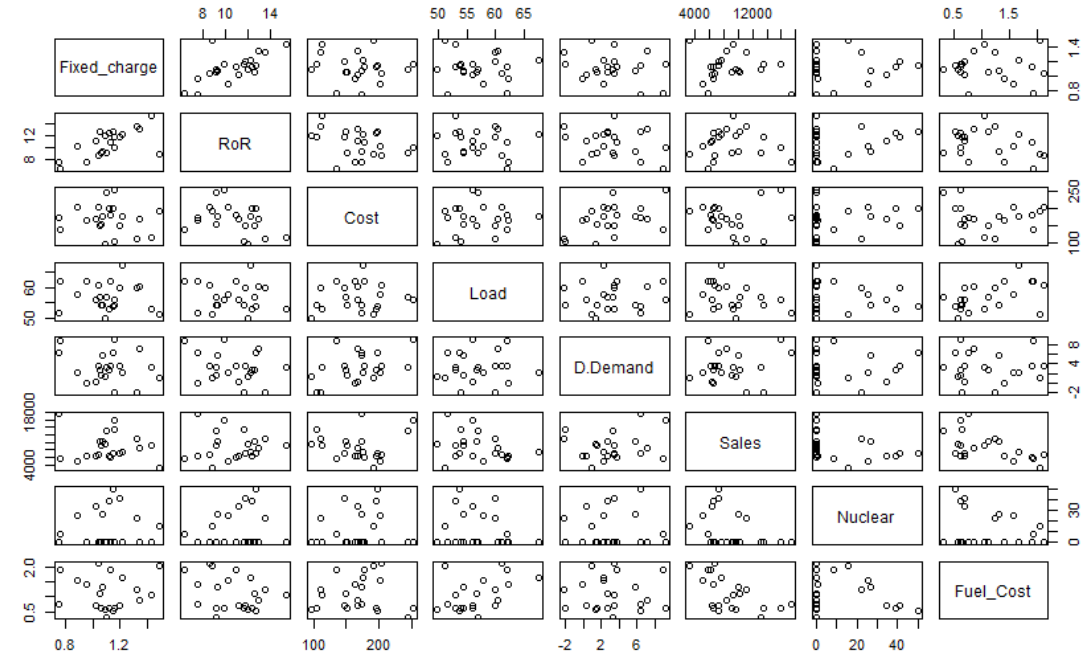
Visualization

```
pairs(mydata[2:9])
```

Scatter plot

```
plot(mydata$Fuel_Cost~ mydata$Sales, data = mydata)  
with(mydata,text(mydata$Fuel_Cost ~ mydata$Sales,  
labels=mydata$Company,pos=3))
```

Scatter plot
between all
possible
combinations
of two
variables



Normalization

```
z <- mydata[,-c(1,1)]
```

```
means <- apply(z,2,mean)
```

```
sds <- apply(z,2,sd)
```

```
nor <- scale(z,center=means,scale=sds)
```

Remove first
column

Normalize

	Fixed_charge	ROR	Cost	Load	D.Demand	Sales	Nuclear	Fuel_Cost
[1,]	-0.29315791	-0.68463896	-0.417122002	-0.57771516	-0.52622751	0.04590290	-0.7146294	-0.85367545
[2,]	-1.21451134	-0.19445367	0.821002037	0.20683629	-0.33381191	-1.07776413	0.7920476	0.81329670
[3,]	1.71214073	2.07822360	-1.339645796	-0.89153574	0.05101929	0.08393124	-0.7146294	-0.08043055
[4,]	-0.50994695	0.20660702	-0.004413989	-0.21906307	-0.94312798	-0.70170610	1.3280197	-0.72420189
[5,]	2.03732429	-0.86288816	0.578232617	-1.29501935	-0.71864311	-1.58142837	0.2143888	1.69263800
[6,]	1.11597086	1.23153991	-1.388199680	0.67756716	-1.74485965	0.62337028	0.6253007	0.24864810
[7,]	0.57399826	0.65223002	0.165524604	2.38116460	-0.33381191	-0.35832428	-0.7146294	0.98772637
[8,]	-0.07636887	-0.68463896	1.864910540	0.00509449	0.01895002	1.17407698	-0.7146294	-1.42731528
[9,]	1.22436538	1.00872841	-0.004413989	0.76723019	1.26965142	-0.14311204	-0.7146294	-0.43288637
[10,]	0.03202565	0.74135462	0.699617327	-0.89153574	-0.17346558	-0.69269198	1.6198267	-0.86266667
[11,]	-1.97327298	-1.44219805	0.116970720	-1.22777208	1.04516655	2.40196983	-0.7146294	-0.60192130
[12,]	0.08622291	0.07292013	0.238355430	1.12588228	0.14722709	-0.77748109	-0.7146294	1.42829614
[13,]	0.19461744	0.87504152	0.748171211	-0.73462545	1.01309729	-0.48874740	2.2749037	-1.03529809
[14,]	-0.13056613	0.56310542	-1.752353809	-1.60883993	-0.59036605	0.21379097	-0.7146294	-0.92560521
[15,]	-0.83513051	-1.39763576	-0.101521757	1.17071379	-1.07140505	-0.68902999	-0.6610322	0.53456889
[16,]	0.24881470	-0.37270287	2.034849134	-0.21906307	1.91103676	1.99351729	-0.7146294	-0.86806140
[17,]	-1.91907572	-1.93238335	-0.781276132	1.10346652	1.84689822	-0.90142531	-0.2203441	1.46965575
[18,]	-0.34735517	0.83047922	-0.441398944	-0.06215278	-0.17346558	0.34534086	-0.7146294	0.00948165
[19,]	0.24881470	0.42941852	-1.558138274	-0.66737818	-1.71279038	1.29379583	-0.7146294	-0.83928950
[20,]	0.46560374	0.47398082	-0.489952828	0.65515141	0.08308855	-0.45832473	1.7329764	-0.72060540
[21,]	-0.40155243	-0.95201276	0.869555920	0.90172472	0.08308855	-0.63776215	-0.7146294	1.82211157

10.3.24 Examples

```
# Calculate distance
distance = dist(nor)
print(distance,digits = 3)
mydata[c(5,11), ]
nor
```

Calculate
Euclidean distance

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21
2	3.10																				
3	3.68	4.92																			
4	2.46	2.16	4.11																		
5	4.12	3.85	4.47	4.13																	
6	3.61	4.22	2.99	3.20	4.60																
7	3.90	3.45	4.22	3.97	4.60	3.35															
8	2.74	3.89	4.99	3.69	5.16	4.91	4.36														
9	3.25	3.96	2.75	3.75	4.49	3.73	2.80	3.59													
10	3.10	2.71	3.93	1.49	4.05	3.83	4.51	3.67	3.57												
11	3.49	4.79	5.90	4.86	6.46	6.00	6.00	3.46	5.18	5.08											
12	3.22	2.43	4.03	3.50	3.60	3.74	1.66	4.06	2.74	3.94	5.21										
13	3.96	3.43	4.39	2.58	4.76	4.55	5.01	4.14	3.66	1.41	5.31	4.50									
14	2.11	4.32	2.74	3.23	4.82	3.47	4.91	4.34	3.82	3.61	4.32	4.34	4.39								
15	2.59	2.50	5.16	3.19	4.26	4.07	2.93	3.85	4.11	4.26	4.74	2.33	5.10	4.24							
16	4.03	4.84	5.26	4.97	5.82	5.84	5.04	2.20	3.63	4.53	3.43	4.62	4.41	5.17	5.18						
17	4.40	3.62	6.36	4.89	5.63	6.10	4.58	5.43	4.90	5.48	4.75	3.50	5.61	5.56	3.40	5.56					
18	1.88	2.90	2.72	2.65	4.34	2.85	2.95	3.24	2.43	3.07	3.95	2.45	3.78	2.30	3.00	3.97	4.43				
19	2.41	4.63	3.18	3.46	5.13	2.58	4.52	4.11	4.11	4.13	4.52	4.41	5.01	1.88	4.03	5.23	6.09	2.47			
20	3.17	3.00	3.73	1.82	4.39	2.91	3.54	4.09	2.95	2.05	5.35	3.43	2.23	3.74	3.78	4.82	4.87	2.92	3.90		
21	3.45	2.32	5.09	3.88	3.64	4.63	2.68	3.98	3.74	4.36	4.88	1.38	4.94	4.93	2.10	4.57	3.10	3.19	4.97	4.15	
22	2.51	2.42	4.11	2.58	3.77	4.03	4.00	3.24	3.21	2.56	3.44	3.00	2.74	3.51	3.35	3.46	3.63	2.55	3.97	2.62	3.01

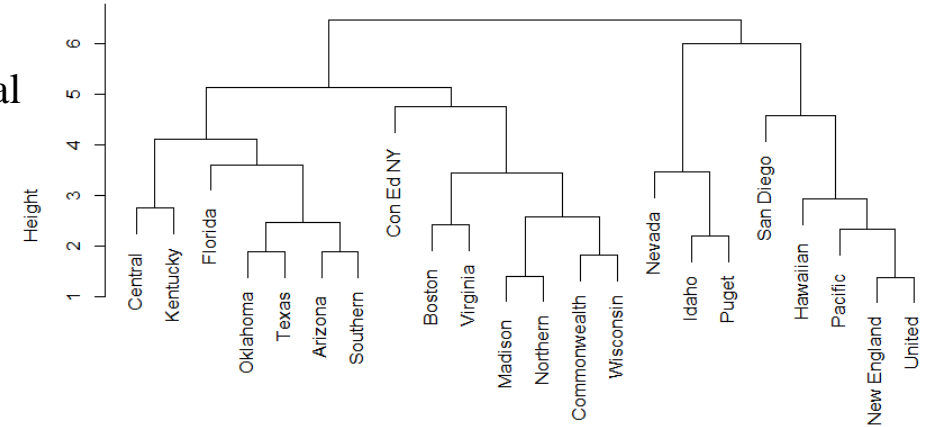
10.3.2.4 Examples

Hierarchical clustering

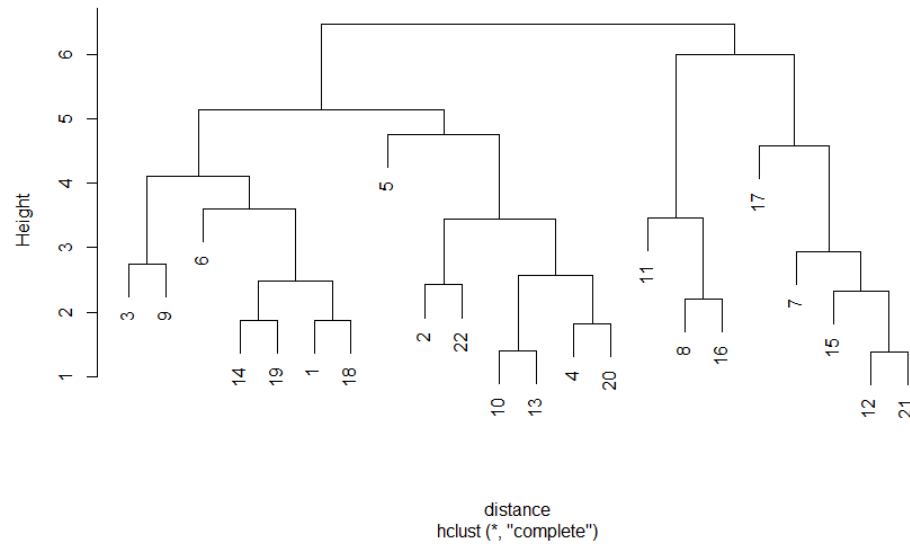
```
mydata.hclust = hclust(distance)
plot(mydata.hclust)
plot(mydata.hclust, labels=mydata$Company, main='Default from hclust')
plot(mydata.hclust, hang=-1)
```

Calculate
Hierarchical
clustering


Default from hclust



Cluster Dendrogram



Cluster Dendrogram



10.3.2.4 Examples

Hierarchical clustering

```
mydata.hclust<-hclust(distance,method="average")  
plot(mydata.hclust,hang=-1)  
rect.hclust(mydata.hclust, k=5)
```

Hierarchical
clustering

Identify clusters from tree

```
member = cutree(mydata.hclust,3)  
table(member)
```


```
> table(member)  
member  
 1  2  3  
18  1  3
```

Characterizing clusters

```
aggregate(nor,list(member),mean)  
aggregate(mydata[,-c(1,1)],list(member),mean)
```

```
> aggregate(nor,list(member),mean)  
Group.1 Fixed_charge RoR Cost Load D.Demand Sales Nuclear Fuel_Cost  
1 1 -0.01313873 0.1868016 -0.2552757 0.1520422 -0.1253617 -0.2215631 0.1071944 0.06692555  
2 2 2.03732429 -0.8628882 0.5782326 -1.2950193 -0.7186431 -1.5814284 0.2143888 1.69263800  
3 3 -0.60027572 -0.8331800 1.3389101 -0.4805802 0.9917178 1.8565214 -0.7146294 -0.96576599  
  
> aggregate(mydata[,-c(1,1)],list(member),mean)  
Group.1 Fixed_charge RoR Cost Load D.Demand Sales Nuclear Fuel_Cost  
1 1 1.111667 11.155556 157.6667 57.65556 2.850000 8127.50 13.8 1.1399444  
2 2 1.490000 8.800000 192.0000 51.20000 1.000000 3300.00 15.6 2.0440000  
3 3 1.003333 8.866667 223.3333 54.83333 6.333333 15504.67 0.0 0.5656667
```

Cluster Dendrogram




10.3.2.4 Examples

Scree plot with sum of square

```
wss <- (nrow(nor)-1)*sum(apply(nor,2,var))
for (i in 2:20)
wss[i] <- sum(kmeans(nor, centers=i)$withinss)
plot(1:20, wss, type="b", xlab="Number of Clusters",
ylab="Within groups sum of squares")
```

WSS to find the
variability, lesser
clusters higher
variability



K-means clustering

```
# K-means clustering
set.seed(123) kc<-kmeans(nor,4)
kc
```

K-Means clustering
with k=4

K-means clustering with 4 clusters of sizes 7, 7, 3, 5

Cluster means:

	Fixed_charge	RoR	Cost	Load	D.Demand	Sales	Nuclear	Fuel_Cost
1	-0.23896065	-0.6591748	0.2556961	0.7992527	-0.05435116	-0.8604593	-0.2884040	1.2497562
2	0.50431607	0.7795509	-0.9858961	-0.3375463	-0.48957692	0.3518600	-0.5232108	-0.4105368
3	-0.60027572	-0.8331800	1.3389101	-0.4805802	0.99171778	1.8565214	-0.7146294	-0.9657660
4	-0.01133215	0.3313815	0.2189339	-0.3580408	0.16646865	-0.4018738	1.5650384	-0.5954476

Clustering vector:

```
[1] 2 1 2 4 1 2 1 3 2 4 3 1 4 2 1 3 1 2 2 4 1 4
```

within cluster sum of squares by cluster:

```
[1] 34.164812 26.507769 9.533522 10.177094
```


```
(between_ss / total_ss = 52.2 %)
```

07-01-2025

Visualization

```
clusplot(mydata,  
         kc$cluster,  
         color = T,  
         shade = T,  
         labels = 2,  
         lines = 0)
```

clusplot uses PCA to
reduce the
dimensions into a
single component



1. Find a decision tree of depth 2 that attains zero training error.
2. Consider the dataset given along. Compute the parameters of Naïve Bayes classifier for predicting **inflated**.

Color	Size	Act	Age	Inflated
YELLOW	SMALL	STRETCH	ADULT	T
YELLOW	SMALL	STRETCH	CHILD	T
YELLOW	SMALL	DIP	ADULT	T
YELLOW	SMALL	DIP	CHILD	T
YELLOW	LARGE	STRETCH	ADULT	T
YELLOW	LARGE	STRETCH	CHILD	F
YELLOW	LARGE	DIP	ADULT	F
YELLOW	LARGE	DIP	CHILD	F
PURPLE	SMALL	STRETCH	ADULT	T
PURPLE	SMALL	STRETCH	CHILD	F
PURPLE	SMALL	DIP	ADULT	F
PURPLE	SMALL	DIP	CHILD	F
PURPLE	LARGE	STRETCH	ADULT	T
PURPLE	LARGE	STRETCH	CHILD	F
PURPLE	LARGE	DIP	ADULT	F
PURPLE	LARGE	DIP	CHILD	F

3. In context of logistic regression, define prediction and loss functions.
4. Discuss the application of Neural network which is used for learning to steer an autonomous vehicle.

The Topics Covered in This Section

10.5.1 Books

10.5.2 Video Lectures

- Tang, P.N., Steibach, M., and Kumar, V. (2016). *Introduction to Data Mining*. Pearson.
 - Provides a theoretical and practical coverage of all data mining techniques with examples. Covers topics like include classification, association analysis, clustering, anomaly detection ★ ★ ★ ★
- Mitchell, T.M. (1997). *Machine Learning*. McGraw Hill.
 - Covers the field of machine learning, which is the study of algorithms that allow computer programs to automatically improve through experience. ★ ★ ★ ★
- Bishop, C.M. (2006). *Pattern Recognition and Machine Learning*. Springer.
 - This book reflects these developments while providing a grounding in the basic concepts of pattern recognition and machine learning. ★ ★ ★
- Hastie, T., Tibshirani, R. Friedman, J. (2009). *The Elements of Statistical Learning*. Springer.
 - Describes the important ideas in a variety of fields such as medicine, biology, finance, and marketing in a common conceptual framework. While the approach is statistical, the emphasis is on concepts rather than mathematics. ★ ★ ★ ★
- James, G., Witten, D., Hastie, T., and Tibshirani, R. (2013). *An Introduction to Statistical Learning*. Springer.

Provides tools for Statistical Learning that are essential for practitioners in science, industry and other fields. Analyses and methods are presented in R. Covers topics like linear regression, classification, resampling methods, shrinkage approaches, tree-based methods, support vector machines, and clustering. ★ ★ ★

- Zaki, M.J., and Jr., W.B. (2014). *Data Mining and Analysis: Fundamental Concepts and Algorithm*. Cambridge University Press.
 - Fundamental algorithms in data mining, automated methods to analyze patterns and models for all kinds of data, with applications ranging from scientific discovery to business intelligence and analytics are covered. Provides a broad yet in-depth overview of data mining, integrating related concepts from machine learning and statistics including exploratory data analysis, pattern mining, clustering, and classification. ★ ★ ★ ★
- Shwartz, S.S., and Davis, S.B. (2014). *Understanding Machine Learning: From Theory to Algorithm*. Cambridge University Press.
 - Introduces machine learning and provides a theoretical account of the fundamentals underlying machine learning and the mathematical derivations that transform these principles into practical algorithms. ★ ★ ★ ★

- Machine Learning | Stanford University

https://www.youtube.com/watch?v=PPLop4L2eGk&list=PLLssT5z_DsK-h9vYZkQkYNWcItqhlRJLN

- Introduction to Machine Learning | University of Toronto

<https://www.youtube.com/watch?v=FvAibtlARQ8&list=PL-Mfq5QSs8iS9XqKuApPE1TSlnZblFHF>

- Machine Learning for Intelligent Systems | Cornell

<https://www.youtube.com/watch?v=MrLPzBxG95I&list=PLl8OlHZGYOQ7bkVbuRthEsaLr7bONzbXS>

- Machine Learning Course | Caltech

<https://www.youtube.com/watch?v=mbyG85GZ0PI&list=PLD63A284B7615313A>

Thank You